



# Introducing Flax NNX:

A Pythonic Neural Network Library for JAX



Helping PyTorch users learn NNX and JAX

# What is Flax NNX?

- Flax NNX
  - A neural network library for JAX, designed for flexibility and high performance
  - New users are encouraged to use Flax NNX
- Flax Linen (Released in 2020)
  - Original neural network library for JAX
  - Focuses on functional programming



# Why Flax NNX?

- Aims to simplify neural network development in JAX.
- **NNX is Pythonic:** Regular Python semantics for Modules, including support for mutability and shared references.
- Emphasizes explicit state management.
- Enables reference sharing and mutability.



# Key Design Principles

- Pythonic Interface
  - Use regular Python objects for defining networks.
- Explicit State Management
  - Deliberate control over parameters and mutable variables.
- Python Graph Data Structure
  - Enables reference sharing and mutability.



# Benefits for PyTorch Users


- Performance
- Smoother transition into the JAX ecosystem.
- API design shares conceptual similarities with PyTorch.
- Familiar patterns for defining neural network architectures.



# Core Concepts

# Core Concept: Modules

- `nnx.Module` is the fundamental building block.
- Modules directly hold their own state (parameters).



# Code Example: Random Layer

```
from flax import nnx
import jax
import jax.numpy as jnp

class RandomLayer(nnx.Module):
    def __init__(self, size: int, *, rngs: nnx.Rngs):
        self.random_vector = nnx.Param(jax.random.normal(rngs.params(),
                                                         (size,)))


    def __call__(self):
        return self.random_vector.value
```



# Flax NNX Python Graphs

## Bridging JAX Power with Python Familiarity



- NNX modules (layers, models, etc.) are regular Python objects that are also registered as JAX Pytrees
- Intuitive, object-oriented feel of Python while enabling seamless integration with JAX's functional transformations like `jax.jit` and `jax.vmap`
- Modules can be composed into nested (graph) structures to express complex models



# Flax NNX Python Graphs

## Bridging JAX Power with Python Familiarity

```
nnx.display(Linear(2, 5, rngs=nnx.Rngs(params=0)))
```

```
▼ Linear( # Param: 15 (60 B)  
  w⇒Param(value=<jax.Array float32(2, 5) ≈0.42 ±0.34 [≥0.029, ≤0.95] nonzero:10>),  
  b⇒Param(value=<jax.Array float32(5,) ≈0.0 ±0.0 [≥0.0, ≤0.0] zero:5>),  
  din=2,   
  dout=5,   
)
```

```
MLP( # Param: 165 (660 B), BatchStat: 32 (128 B), RngState: 2 (12 B), Total: 199 (800 B)
  linear1=Linear( # Param: 48 (192 B)
    w=Param(value=<jax.Array float32(2, 16) ≈0.51 ±0.33 [≥0.0059, ≤0.97] nonzero:32>),
    b=Param(value=<jax.Array float32(16,) ≈0.0 ±0.0 [≥0.0, ≤0.0] zero:16>),
    din=2,
    dout=16,
  ),
  dropout=Dropout( # RngState: 2 (12 B)
    rate=0.1,
    broadcast_dims=(),
    deterministic=False,
    rng_collection='dropout',
    rngs=Rngs( # RngState: 2 (12 B)
      default=RngStream(key=RngKey(value=<jax.Array key<fry>()>, tag='default'), count=R
    ),
  ),
  bn=BatchNorm(mean=BatchStat(value=<jax.Array float32(16,) ≈0.01 ±0.0051 [≥0.0013, ≤0.0
  linear2=Linear(w=Param(value=<jax.Array float32(16, 5) ≈0.51 ±0.3 [≥0.0061, ≤1.0] nonz
)
```

# Data Structure: Python Graphs

- **Direct Attribute Access:** Access layers and parameters like `model.layer1`, `model.layer1.weight`
- **Standard Python References:** Assigning a layer to multiple attributes means they share the exact same object, just like in Python. Useful for weight sharing!
- **Mutability:** Module state can be modified directly



# Python Graphs in Practice: Like `torch.nn.Module` but for JAX


## Handling State:

- **Static Config:** Regular Python attributes
  - e.g., `self.dropout_rate = 0.5`
- **Dynamic State (Parameters/Buffers):** Stored in special `nnx.Variable` container objects
  - e.g., `nnx.Param` for trainable weights
- Access the actual JAX array data via `.value`
  - e.g., `self.weight.value`

# Flax NNX Python Graphs

## Bridging JAX Power with Python Familiarity

- Why Python Graphs?
  - **Mutability:** Allowing objects (like model layers) to change their internal state directly.
  - **Reference Sharing:** Ensuring that if you assign the same layer object to multiple places in your model, it's truly the same object instance.



# Python Graphs in Practice: Like `torch.nn.Module` but for JAX

## Reference Sharing:

- Assigning the same module or `nnx.Variable` instance to multiple attributes creates shared references, identical to Python's standard behavior

```
shared_layer = nnx.Linear(...)  
model.encoder_layer = shared_layer  
model.decoder_layer = shared_layer # Both point to the same layer.
```

# Variable Types

**Variables** are wrappers over `jax.Arrays` that enable state updates and allow associated metadata.

- `nnx.Param`: Learnable parameters of the model (dynamic state).
- `nnx.BatchStat`, `nnx.Cache`, `nnx.Intermediate`: For specialized state.
- Also `nnx.State` (not a **Variable** type): A pytree that contains a pure version of the object's state.





# Code Example: Stateful Parameter Layer

```
from flax import nnx
import jax.numpy as jnp

class StatefulParameterLayer(nnx.Module):
    def __init__(self, initial_value: float, *, rngs: nnx.Rngs):
        self.weight = nnx.Param(jax.random.uniform(rngs.params()))
        self.bias = {'bias': jnp.array(initial_value)}

    def update_bias(self, new_value: float):
        self.bias['bias'] = jnp.array(new_value)

    def __call__(self, x: jax.Array):
        return x * self.weight.value + self.bias['bias']
```


# Explicit Random Number Generation

- Uses `nnx.Rngs` object.
- Requires explicit passing of a PRNG key when instantiating modules with randomness.
- Promotes reproducibility and easier parallelization.
- Layers store a **forked copy** of RNGs, ensuring state isolation.



# Eager Parameter Initialization

- Parameters are initialized immediately when an `nnx.Module` is instantiated.
- All shape information must be provided during initialization.
- No implicit shape inference.



# Functional API: split, merge, update

- `nnx.split`: Decomposes an `nnx.Module` into its static structure (`GraphDef`) and dynamic state (`State`).
- `nnx.merge`: Reconstructs an `nnx.Module` from its `GraphDef` and `State`.
- `nnx.update`: Updates an existing object in-place with the content of a given `State`.



# Code Example: Counter with Functional API

```
class Counter(nnx.Module):
    def __init__(self):
        self.count = 0

    def __call__(self):
        self.count += 1
        return self.count

counter = Counter()
print(f'{counter() = }') # 1
graphdef, state = nnx.split(counter)

@jax.jit
def count(state):
    counter = nnx.merge(graphdef, state)
    counter() # 2
    return nnx.state(counter)

state = count(state)
nnx.update(counter, state)
print(f'{counter() = }') # 3
```

# nnx.jit: When to Use NNX's Supercharger

- **What it does:** `nnx.jit` ("Just-In-Time" compilation) takes your Python function and uses XLA to compile it into highly optimized code for CPU/GPU/TPU.
- **Think:** Similar goal to `torch.compile()` – speed up execution.
- **The Catch:** You shouldn't just `jit()` every function! Performance gains come from JITting *the right* functions.
- **Key Requirement #1:** `nnx.jit` works best on Pure Functions. We'll define this next.
- **Goal:** Identify the computationally heavy, pure parts of your code and `jit()` those.

# nnx.jit: NNX's Supercharger

```
import jax
import jax.numpy as jnp
from flax import nnx
import torch

jax_w = jnp.ones((3, 4))
jax_b = jnp.ones(3)
jax_x = jnp.ones(4)

pt_w = torch.ones((3, 4))
pt_b = torch.ones(3)
pt_x = torch.ones(4)

def jax_predict(W, b, x):
    return jnp.matmul(W, x) + b

def pt_predict(W, b, x):
    return torch.matmul(W, x) + b
```

# nnx.jit: NNX's Supercharger

```
jit_predict = nnx.jit(jax_predict)

def jax_batch(W, b, x):
    for i in range(10000):
        result = jax_predict(W, b, x)

def pt_batch(W, b, x):
    for i in range(10000):
        result = pt_predict(W, b, x)

%timeit pt_batch(pt_w, pt_b, pt_x) # PyTorch
%timeit jit_predict(jax_w, jax_b, jax_x).block_until_ready() # JAX with XLA compile
%timeit jit_predict(jax_w, jax_b, jax_x).block_until_ready() # NNX with jit()
```

79.9 ms  $\pm$  12.5 ms per loop (mean  $\pm$  std. dev. of 7 runs, 10 loops each)  
The slowest run took 6.12 times longer than the fastest. This could mean that an intermediate result is being cached.

**79.4  $\mu$ s**  $\pm$  72.3  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

**39.9  $\mu$ s**  $\pm$  756 ns per loop (mean  $\pm$  std. dev. of 7 runs, 10000 loops each)



# nnx.jit: What is a "Pure Function"? (The Key to JIT)

Definition: A function is pure iff:

- It always returns the same output for the same inputs
- It has no side effects (we'll talk about this next)
- **The Rule: Only apply `nnx.jit` to functions that are pure**
- (Good Practice: Writing pure functions is often good design anyway – easier to test and debug!)\*

# nnx.jit: What is a "Pure Function"? (The Key to JIT)

## What are Side Effects?

- Actions that affect things outside the function's direct input/output:
  - Modifying global variables
  - Changing mutable inputs (e.g., appending to a list passed as an argument)
  - Printing to the console, writing to files (I/O)
  - Changing system settings, interacting with databases

## Another Reason Not to jit: Python Control Flow

**The Scenario:** Your function uses standard Python `if`, `while`, `for` loops where the condition or iteration *depends directly on the value* of an input argument.

```
# Example NOT ideal for direct jit()

def process_value(x, threshold):
    if x > threshold: # Control flow depends on input values
        return x * 2
    else:
        return x / 2
```

# Another Reason Not to jit: Python Control Flow

Why it's Problematic for `jit()`:

- JAX traces the function based on the shapes and types of inputs, not specific values (usually)
- A standard Python `if` depending on a value creates a specific path during tracing
- If you call the JITted function later with a value that takes a different path, JAX might error or need to recompile, which can be slow, losing the speed benefit

# Another Reason Not to jit: Python Control Flow

**The Rule: Avoid JITting functions where standard Python control flow (if/while) depends directly on input values.**


Strategy:

- Isolate the pure, numerically intensive parts within the branches/loops into separate functions
- `jit()` those smaller, pure functions
- Keep the Python control flow logic outside the JITted functions

# JAX Transformations & The Stateful Model Challenge

Standard JAX Transformations  
(`jax.jit`, `jax.grad`, `jax.vmap`)

- Designed for **pure functions**:
  - No side effects (don't modify external state).
  - Deterministic output for the same input.
- Operate on **PyTrees** (nested lists, dicts, tuples containing arrays).
- Require **explicit state management**: You must manually pass state (like model parameters, optimizer state, RNG keys) into the function and return the updated state.



# JAX Transformations & The Stateful Model Challenge


NNX Modules (and PyTorch Modules) are stateful:

- NNX Modules (like `torch.nn.Module`) are inherently stateful. They contain their parameters and potentially other mutable state (e.g., BatchNorm stats).
- Directly applying `jax.jit()` or `jax.grad()` to a method of a stateful object requires manually extracting the state, passing it through the pure function, and then updating the object with the returned state. This can be cumbersome.



# NNX Transformations

- `nnx.jit`, `nnx.grad`, `nnx.vmap`
- Wrappers around standard JAX transformations
  - `jax.jit`, `jax.grad`, `jax.vmap`, etc.
- Specifically designed to work directly with NNX graph objects
  - `nnx.Module`, `nnx.Optimizer`, `nnx.Rngs`, etc.





# The Key Difference: Automatic State Management

**JAX Transforms:** Require YOU to handle state explicitly (pass in, get back).

```
# Simplified JAX pattern
params, opt_state = ...
grads, new_state = jax.grad(loss_fn, has_aux=True)(params, ...)
updates, opt_state = optimizer.update(grads, opt_state, params) # Pass opt_state in, get it back
params = apply_updates(params, updates)
```

**NNX Transforms:** Handle state lifting and updating automatically behind the scenes when applied to methods of NNX objects.

```
# Simplified NNX pattern
model, optimizer, rngs = ... # NNX objects holding state
@nnx.grad # Applied to a method or function working with NNX objects
def loss_fn_nnx(...): ...
grads = loss_fn_nnx(...)
optimizer.update(grads) # Updates parameters within model object directly
```

# When to Use NNX vs. JAX Transformations

Use NNX Transformations (`nnx.jit`, `nnx.grad`, etc.) when:



- You are working with NNX graph objects (`nnx.Module`, `nnx.Optimizer`, `nnx.Variable`, `nnx.Rngs`). This is the primary use case.
- You want simplified state management and less boilerplate code.
- You prefer a more object-oriented style where transformations apply directly to methods interacting with stateful objects.



# When to Use NNX vs. JAX Transformations

Use Standard JAX Transformations (`jax.jit`, `jax.grad`, etc.) when:

- You are working with pure functions that don't involve NNX objects or mutable state directly.
- You are operating on plain **PyTrees** (e.g., data preprocessing functions).
- You need fine-grained, low-level control over the transformation process and state handling (you want to manage state explicitly).
- You need a specific JAX transformation feature that might not yet have an NNX counterpart (less common for core transforms).
- Standard JAX transformations are faster



# When to Use NNX vs. JAX Transformations

General Recommendation:


- **If you are using Flax NNX objects, use the corresponding NNX transformations.** They are designed for this purpose and provide a much smoother experience.



# Models


# Fundamental Neural Network Layers

- `nnx.Linear`
- `nnx.Conv`
- `nnx.BatchNorm`
- `nnx.LayerNorm`
- `nnx.GroupNorm`
- `nnx.MultiHeadAttention`
- `nnx.LSTMCell`
- `nnx.GRUCell`
- `nnx.Dropout`




# Building Complex Models

- MLPs, CNNs can be easily constructed.
- Requires explicit specification of input and output shapes during initialization.
- Provision of an appropriate random number generator key for parameters.



# MNIST Tutorial

- Example of defining a CNN for digit classification using Flax NNX.
- Covers loading the MNIST dataset, defining the CNN model, creating an optimizer using Optax, defining the training loop, and evaluating the model.



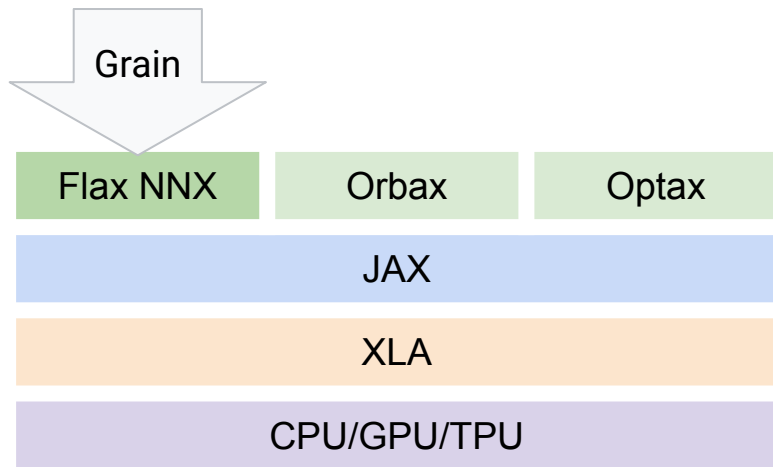
[https://flax.readthedocs.io/en/latest/mnist\\_tutorial.html](https://flax.readthedocs.io/en/latest/mnist_tutorial.html)



# Flax NNX and the JAX AI Stack

# Flax NNX and the JAX AI Stack


- **JAX**: Array operations and program transformations.
- **Flax NNX**: Building neural networks.
- **Optax**: Gradient processing and optimization.
- **Orbax**: Checkpointing and persistence.
- **ml\_dtypes**: NumPy dtype extensions for machine learning.



# Comparing Flax NNX with PyTorch


# Comparison with PyTorch: High-Level

- Both use class-based structures for defining models.
- Similarities in `__init__` and forward pass methods (`__call__` in NNX).
- Key difference: explicit vs. implicit random number generation.



# Comparison: State Management

- **PyTorch:** Imperative updates using `zero_grad()` and `step()`.
- **Flax NNX:** Direct attribute assignment and mutation, also enables a functional style with Functional API (`split`, `merge`).



# Code Example: Shifted ReLU in PyTorch and NNX

```
# PyTorch
import torch
import torch.nn as nn
```

```
class ShiftedReLU_Torch(nn.Module):
    def __init__(self, shift: float):
        super().__init__()
        self.shift = shift

    def forward(self, x):
        return torch.relu(x + self.shift)
```

```
# Flax NNX
from flax import nnx
import jax.numpy as jnp
```

```
class ShiftedReLU_NNX(nnx.Module):
    def __init__(self, shift: float):
        self.shift = shift

    def __call__(self, x: jnp.ndarray):
        return nnx.relu(x + self.shift)
```

# Code Example: Simple Classifier in PyTorch and NNX

```
# PyTorch
import torch
import torch.nn as nn

class SimpleClassifier_Torch(nn.Module):
    def __init__(self, input_size: int,
                  num_classes: int):
        super().__init__()
        self.linear1 = nn.Linear(
            input_size, 10)
        self.relu = nn.ReLU()
        self.linear2 = nn.Linear(
            10, num_classes)

    def forward(self, x):
        x = self.linear1(x)
        x = self.relu(x)
        return self.linear2(x)
```

```
# Flax NNX
from flax import nnx
import jax
import jax.numpy as jnp

class SimpleClassifier_NNX(nnx.Module):
    def __init__(self, input_size: int,
                  num_classes: int, *,
                  rngs: nnx.Rngs):
        self.linear1 = nnx.Linear(
            input_size, 10, rngs=rngs)
        self.relu = nnx.relu
        self.linear2 = nnx.Linear(
            10, num_classes, rngs=rngs)

    def __call__(self, x: jnp.ndarray):
        x = self.linear1(x)
        x = self.relu(x)
        return self.linear2(x)
```

# Comparison: Training Loop & Backpropagation

- **PyTorch:** `loss.backward()` for automatic gradient computation, optimizer updates parameters directly.
- **Flax NNX:** `nnx.value_and_grad` to compute gradients and `optimizer.update` to update the model's state with computed gradients.





# Comparison: Training Loop & Backpropagation

```
# PyTorch
import torch
import torch.nn as nn
import torch.optim as optim
```

```
# Define a simple model
class SimpleModel_Torch(nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = nn.Linear(1, 1)

    def forward(self, x):
        return self.linear(x)
```

```
# Flax NNX
from flax import nnx
import jax
import jax.numpy as jnp
from optax import sgd
from typing import Any
```

```
# Define a simple model
class SimpleModel_NNX(nnx.Module):
    def __init__(self, *, rngs: nnx.Rngs):
        self.linear = nnx.Linear(1, 1, rngs=rngs)

    def __call__(self, x: jnp.ndarray):
        return self.linear(x)
```

# Comparison: Training Loop & Backpropagation

# PyTorch

```
model_torch = SimpleModel_Torch()
```

```
optimizer_torch = \
    optim.SGD(model_torch.parameters(), lr=0.01)
loss_fn = nn.MSELoss()
```

# Dummy data

```
x_torch = torch.tensor(
    [[2.0]], requires_grad=True)
y_torch = torch.tensor([[4.0]])
```

# Flax NNX

```
model_nnx = SimpleModel_NNX(rngs=nnx.Rngs(0))
```

# Optimizer

```
optimizer = nnx.Optimizer(
    model_nnx,
    tx=sgd(learning_rate=0.01),
    wrt=nnx.Param)
```

# Dummy data

```
x_nnx = jnp.array([[2.0]])
y_nnx = jnp.array([[4.0]])
```

# Comparison: Training Loop & Backpropagation

```
# PyTorch Training step

# Zero the gradients
optimizer_torch.zero_grad()
output_torch = model_torch(x_torch)
loss_torch = loss_fn(
    output_torch, y_torch)
loss_torch.backward() # Compute gradients

# Update parameters
optimizer_torch.step()

print("PyTorch Loss:", loss_torch.item())
```

```
# Flax NNX Training step
@nnx.jit
def train_step(model, optimizer, x, y):
    def loss_fn(model):
        return jnp.mean((model(x) - y) ** 2)

    loss, grads = \
        nnx.value_and_grad(loss_fn)(model)
    # in-place updates
    optimizer.update(model, grads)

    return loss

# Pass the optimizer
loss_nnx = train_step(model_nnx,
                       optimizer, x_nnx, y_nnx)

print("Flax NNX Loss:", loss_nnx)
```

# Conclusion

- Flax NNX provides a powerful and intuitive way to build neural networks with JAX.
- **NNX is Pythonic:** Regular Python semantics for Modules, including support for mutability and shared references.



# Learning Resources

Code Exercises, Quick References, and Slides

- <https://goo.gle/learning-jax>



# Community and Docs

## Community:

- <https://goo.gle/jax-community>

## Docs

- JAX AI Stack: <https://jaxstack.ai>
- JAX: <https://jax.dev>
- Flax NNX: <https://flax.readthedocs.io>