



Debugging JAX & Flax NNX:

A Guide for PyTorch Users



Leveraging Familiar Concepts and Mastering New Tools

Why is JAX Debugging Different? The JIT Impact

- **PyTorch:** Mostly eager execution. `print()`, `pdb` work directly on runtime values.
- **JAX:** Relies heavily on Just-In-Time (JIT) compilation (`@jax.jit`) for performance.
 - **Tracing Phase:** Python code runs once with abstract tracers (shapes/types) to build a computation graph. Standard `print()` or `pdb` only see these tracers. However tracers can be useful, since they carry shapes and `dtypes`.
 - **Execution Phase:** Compiled graph (e.g., XLA) runs later on CPU/GPU/TPU with concrete values. Standard tools can't inspect this directly.
- **Challenge:** How to inspect *runtime* values inside JIT-compiled code?

"printf Debugging" in JAX: `jax.debug.print()`

- The JAX equivalent of `print()` for use inside transformed functions (`jit`, `vmap`, `grad`).
- Embeds the print operation into the compiled computation graph.
- Outputs concrete runtime values during execution.
- Use `ordered=True` to ensure prints appear sequentially as written.

"printf Debugging" in JAX: `jax.debug.print()`

```
@jax.jit
def compute_intermediate(x):
    y = x * 2
    print("Standard print (tracer):", y)    # Sees tracers during compilation
    # Sees runtime values
    jax.debug.print("jax.debug.print (runtime value): {y}", y=y, ordered=True)
    z = y + 1
    return z
```

```
compute_intermediate(jnp.array(5.0))
```

Output:

```
# Standard print (tracer): Traced<ShapedArray(float32)...>
```

```
# jax.debug.print (runtime value): 10.0
```

Interactive Debugging in JIT: `jax.debug.breakpoint()`

- The JAX equivalent of `pdb.set_trace()` or `breakpoint()` for use inside transformed functions.
- Pauses execution at runtime within the compiled code.
- Provides a (`jaxdb`) prompt similar to `pdb`.
- Allows inspecting runtime values of variables in the JAX context (`p variable_name`).
- Use `c` to continue, `q` to quit.
- Can be made conditional using `jax.lax.cond()`.

Interactive Debugging in JIT: `jax.debug.breakpoint()`

```
@jax.jit
def check_value(x):
    y = jnp.sin(x)
    jax.debug.print("Value before breakpoint: {y}", y=y)
    jax.debug.breakpoint() # Execution pauses here
    z = jnp.cos(y)
    return z
```

```
check_value(jnp.array(0.5))
```

```
# Output: Value before breakpoint: 0.47942555
```

```
# (jaxdb) p y
```

```
# Array(0.47942555, dtype=float32)
```

```
# (jaxdb) c
```

Visualizing Data Layout: `jax.debug.visualize_array_sharding`


- **Problem:** Need to verify data layout across devices in distributed settings with sharding.
- **Tool:** `jax.debug.visualize_array_sharding(array)` prints text diagram of layout.
- **How it Works:** Shows which data slice is on which device ID at runtime.
- **Usage:** Call inside JIT/distributed functions to check sharding from `Mesh/PartitionSpec`.

visualize_array_sharding Code Example


```
# --- Assume mesh setup (e.g., 2x2) ---  
# mesh = Mesh(devices, ('data', 'model'))  
# P = PartitionSpec  
  
@jax.jit  
def my_distributed_func(x_sharded):  
    # Visualize input array layout  
    print("--- Input Sharding ---")  
    jax.debug.visualize_array_sharding(x_sharded) # <<< Visualize  
  
    y = x_sharded * 2 # Computation  
    print("--- Output Sharding ---")  
    jax.debug.visualize_array_sharding(y) # <<< Visualize  
    return y
```


Visualizing Data Layout: `jax.debug.visualize_array_sharding`

`sharded_model.dot1.kernel (None, 'model') :`



`sharded_model.w2 ('model', None) :`



Back to Basics: Temporarily Disabling JIT

Sometimes you just want standard Python debugging.

- `jax.disable_jit()` forces JAX functions to execute eagerly (like PyTorch/NumPy).
- Allows standard `print()` and `pdb/breakpoint()` to work as expected, inspecting runtime values directly.
- Also enables IDE debuggers like VS Code

Back to Basics: Temporarily Disabling JIT

- How to use:
 - **Context Manager:** `with jax.disable_jit():...` (Recommended for locality)
 - **Globally:** `jax.config.update("jax_disable_jit", True)`
 - **Environment Var:** `JAX_DISABLE_JIT=1`
- **Major Drawback:** Disables JIT optimizations -> significantly slower execution. Use temporarily for debugging!

Back to Basics: Temporarily Disabling JIT

```
@jax.jit
def problematic_function(x):
    y = jnp.log(x)
    # pdb works here ONLY if JIT is disabled
    pdb.set_trace()
    return y * 2

# Execute with JIT disabled for this block
with jax.disable_jit():
    print("Running with JIT disabled...")
    result_no_jit = problematic_function(jnp.array(5.0)) # pdb triggers
```

Back to Basics: Temporarily Disabling JIT

Limits on `jax.disable_jit()`:

- If you're using functional transforms, like `jax.vmap` and `jax.scan`, you won't be able to break into the function to inspect values
- But tools like `sow` from NNX are designed to be compatible with these transforms

Automatic NaN Hunting: The `jax_debug_nans` Flag

Problem: Numerical instability leading to NaNs inside JITted code can be hard to trace.

Solution: Enable the `jax_debug_nans` configuration flag.

- `jax.config.update("jax_debug_nans", True)` (in Python)
- Environment Variable: `JAX_DEBUG_NANS=1`

Automatic NaN Hunting: The `jax_debug_nans` Flag

How it Works:

- Monitors JIT computations for **NaN** outputs.
- If a **NaN** is detected, JAX automatically re-runs the function in eager mode (like `disable_jit`) to pinpoint the exact operation causing the **NaN** and raise an error there.

Automatic NaN Hunting: The `jax_debug_nans` Flag

Limitations:

- Can significantly slow down execution due to checks and potential eager re-runs.
- Might raise errors on intentionally created **NaNs**.
- Best used during debugging, disable for production/performance runs.

Inspecting Flax NNX Models: `nnx.display()`

Understanding Your NNX Model: `nnx.display()`

Flax NNX is designed with inspectability in mind (uses standard Python objects)

- `nnx.display()` provides a clear view of NNX objects:
 - `nnx.Module` (models, layers)
 - `nnx.Optimizer`
 - `nnx.State`
 - Other contained JAX arrays/objects

Understanding Your NNX Model: `nnx.display()`

- Shows structure, parameters, state variables, shapes, types, and values.
- Creates a rich, interactive tree view in notebooks/Colab. Falls back to standard print otherwise.
- Analogous to `print(pytorch_model)` but often more detailed and interactive.

```
▼ Linear( # Param: 15 (60 B)
  w⇒Param(value=<jax.Array float32(2, 5) ≈0.42 ±0.34 [≥0.029, ≤0.95] nonzero:10>),
  b⇒Param(value=<jax.Array float32(5,) ≈0.0 ±0.0 [≥0.0, ≤0.0] zero:5>),
  din=2,
  dout=5,
)
```

Capturing Intermediate Values: `nnx.sow()`

Flax NNX Module.sow() - Capturing Intermediate Values

What is Module.sow()?

- A method within `flax.nnx.Module` designed to capture and store intermediate values computed during a module's forward pass.
- **Purpose:** Simplifies tracking internal computations (e.g., activations, gradients) without manually passing data structures through the call stack.
- Useful for:
 - Debugging model behavior.
 - Visualizing intermediate layer outputs.
 - Implementing custom loss functions based on internal states.

Flax NNX Module.sow() - Capturing Intermediate Values

How it Works: Core Arguments

```
self.sow(variable_type, name, value, *, reduce_fn=None, init_fn=None)
```

- **variable_type**: Specifies the variable wrapper (e.g., `nnx.Intermediate`). Used for later filtering/retrieval.
- **name (str)**: The attribute name under which the value will be stored on the module instance.
- **value**: The actual JAX array or Python object to be stored.
- **(Optional) reduce_fn, init_fn**: Customize how values are stored/aggregated if `sow` is called multiple times with the same name

Using `Module.sow()` - Storage, Retrieval & Customization

```
class SimpleModel(nnx.Module):  
    def __init__(self, rngs):  
        self.dense = nnx.Linear(2, 3, rngs=rngs)  
    def __call__(self, x):  
        x = self.dense(x)  
        # Sow the output of the dense layer  
        self.sow(nnx.Intermediate, 'dense_output', x)  
        return x * 2
```

```
model = SimpleModel(rngs=nnx.Rngs(0))
```

```
y = model(jnp.ones((1, 2)))
```

```
# Retrieve the sown value
```

```
intermediate_val = model.dense_output.value
```

```
# Output: (array([[...]]),) - A tuple containing the value
```

```
print(intermediate_val)
```

Flax NNX Module `.sow()` - Capturing Intermediate Values

Storage and Retrieval

- Sown values are stored as attributes on the module instance, named according to the name argument.
- The actual data is accessed via the `.value` property of this attribute (e.g., `model.dense_output.value`).
- **Default Behavior:** Calling `sow` multiple times with the same name appends the new value to a tuple stored in `.value`. Be mindful of potential memory growth with frequent calls.

Flax NNX Module.sow() - Capturing Intermediate Values

Advanced Features

- **Custom Aggregation:** Use `init_fn` and `reduce_fn` to define custom logic for combining multiple sown values (e.g., sum, product, average) instead of appending to a tuple.
- **State Management:** Sown variables (typed by `variable_type`) integrate with NNX graph utilities like `nnx.split()`, `nnx.state()`, and `nnx.pop()` for extracting or managing specific types of intermediate values from the module state.

Robustness with Chex Assertions

Note:

This section is optional if you have already gone through the Chex module.

Catching Errors Early: Chex Assertions

Chex: A library from DeepMind for reliable JAX code (testing, debugging)

- Provides powerful assertion functions.
- Static Assertions: Check properties independent of runtime values (shapes, dtypes, rank, structure).
- Work seamlessly inside `@jax.jit()` because they operate on traced information.

Catching Errors Early: Chex Assertions

Examples (static assertions):

- `chex.assert_shape(array, (batch, features))`
- `chex.assert_rank(array, 2)`
- `chex.assert_type(array, jnp.float32)`
- `chex.assert_trees_all_equal_shapes(tree1, tree2)`

Monitoring with TensorBoard

Visualizing Training: TensorBoard Integration (Setup)

TensorBoard works well with JAX/Flax for monitoring experiments.

Setup:

1. **Install:** `pip install tensorboard` (You might need tensorflow or torch[tensorboard] for the writer).
2. **Create Summary Writer:** Points to a log directory. Can use TensorBoardX, or TensorFlow, or PyTorch utilities.
3. **Launch TensorBoard:** From terminal: `tensorboard --logdir logs`
4. **Access in browser** (usually `http://localhost:6006`).

TensorBoard: Create Summary Writer

Option 1: TensorBoardX

```
from tensorboardX import SummaryWriter  
writer = SummaryWriter("logs/my_run_1")
```

Option 2: TensorFlow writer

```
import tensorflow as tf  
writer = tf.summary.create_file_writer("logs/my_run_1")
```

Option 3: PyTorch writer (if torch installed)

```
from torch.utils.tensorboard import SummaryWriter  
writer = SummaryWriter("logs/my_run_1")
```

Visualizing Training: TensorBoard Integration (Logging)

- Log metrics, images, etc., from your training loop.
- **Scalars** (Loss, Accuracy):
 - Use `tf.summary.scalar(...)` (TF writer, or PyTorch writer)
 - **Crucial:** Convert JAX arrays to Python scalars using `.item()` before logging.
- **Images:** Visualize inputs, predictions, attention maps (`writer.add_image(...)`).
- **Profiling:** JAX profiling (`jax.profiler`) can also output data viewable in TensorBoard for performance analysis.

Code Example (Logging Scalars with TensorBoardX Writer)

```
# Inside training loop (epoch = current step)
# Assume train_loss, val_accuracy are JAX arrays

writer.add_scalar('Loss/train', train_loss.item(), global_step=epoch)
writer.add_scalar('Accuracy/validation', val_accuracy.item(), global_step=epoch)

# Remember writer.close() when done
```

- ☐ Show data download links
- ☒ Ignore outliers in chart scaling

Tooltip sorting method: **default** ▼

Smoothing

 0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs



TOGGLE ALL RUNS

logs/jax_debug_run_solution

🔍 Filter tags (regular expressions supported)


Accuracy

Accuracy/validation
tag: Accuracy/validation



Loss

Loss/train
tag: Loss/train



Profiling with XProf, aka JAX Profiler

Installing and Enabling XProf

Installing and enabling XProf in TensorBoard is easy!

XProf includes:

- Overview Page, Framework Op Stats, Graph Viewer, HLO Op Stats, Memory Profile, Memory Viewer, HLO Op Profile, Roofline Model, Trace Viewer

```
!pip install -Uq tensorboard tensorboard_plugin_profile  
jax.profiler.start_trace(LOG_DIR) # Capturing trace for xprof  
  
// Some training loop code  
  
jax.profiler.stop_trace()
```

≡ XProf

CAPTURE PROFILE

Runs (1)

tpu-training ▼

Tools (9)

Overview Page ▼

Hosts (1)

gke-tpu-b309f56b-rq5s ▼

Performance Summary

Average Step Time **6676.2 ms**
($\sigma = 45.5$ ms)
lower is better.

FLOPS Utilization ⓘ
higher is better.


Utilization of TPU
Matrix Units 30.6%

Compared to
Program's Optimal
FLOPS 0.0%
see [roofline_model](#)

TPU Duty Cycle ⓘ 0.0%
higher is better.

**Memory Bandwidth
Utilization** ⓘ 0.0%
higher is better.

Step-time Graph



Digging into your model with Model Explorer

Exploring the Model with Model Explorer

Model Explorer is a powerful graph visualization tool designed to help you understand and debug your ML models.

- Hierarchical Layout
 - View your model's structure clearly with nested layers that you can expand and collapse as needed.
- Metadata Overlays
 - Gain insights by overlaying metadata (e.g., attributes, inputs/outputs, etc) and custom data (e.g. performance) on nodes.
- Powerful Interactive Features
 - Focus on specific areas with search, navigate graphs smoothly with bookmarks and layer popups, customize node styles with queries, compare graphs side-by-side, and more.

Search nodes/layers by regex



dd

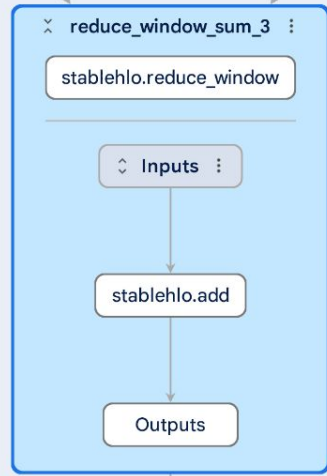
stablehlo.convert

stablehlo.concatenate

func.call

stablehlo.broadcast_in_dim

stablehlo.dynamic_broadcast_in_dim



stablehlo.reshape

stablehlo.divide

stablehlo.reshape

stablehlo.concatenate

stablehlo.convolution

stablehlo.dynamic_broadcast_in_dim

LAYER INFO

name reduce_window_sum_3
namespace jit(forward_jax)/jit(main)
#children 4
#descendants 6

IDENTICAL LAYERS (2)

<root> / jit(forward_jax) / jit(main)
reduce_window_sum_3
reduce_window_sum_15

Op
Layer
Selected layer
Identical layer (if any)

Zoom
Pan
Ctrl+Scroll
Drag or scroll

Bridging from PyTorch:

Similarities, Key Differences, & Adaptations

Mapping Your PyTorch Skills to JAX/NNX

Your PyTorch Toolkit:

- `print()` for values
- `pdb` / `breakpoint()` for interactive steps
- Manual NaN/Inf checks
- `print(model)` for structure
- TensorBoard
- Hooks for intermediate values
- Profiler

Mapping Your PyTorch Skills to JAX/NNX

JAX/Flax NNX Analogues:

- **Value Inspection:** `jax.debug.print()` (in JIT), `print()` (if using `jax.disable_jit`)
- **Interactive Debugging:** `jax.debug.breakpoint` (in JIT), `pdb` or IDE (if using `jax.disable_jit`)
- **NaN/Inf Checks:** `chex.assert_tree_all_finite` (with `chex.chexify`)
- **Structure Inspection:** `nnx.display(model)`
- **Monitoring:** TensorBoard (similar setup/usage)
- **Assertions:** Chex (`assert_shape`, etc.)

Key Differences & Necessary Adaptations

- JIT is King: The biggest change. Debugging often means choosing:
 - Use JAX-specific tools (`jax.debug.*`, Chex + `chex.chexify`) to work within JIT.
 - Temporarily disable JIT (`jax.disable_jit`) to use standard tools, sacrificing performance.

Key Differences & Necessary Adaptations

- Hooks vs. Functional Style: PyTorch uses hooks heavily. JAX/NNX leans towards:
 - Returning intermediate values explicitly from functions (e.g., using `has_aux=True` in `jax.grad`).
 - Using `jax.debug.callback` for more complex host interactions.
 - Transforming functions (like gradient transforms) instead of in-place modification.

Key Differences & Necessary Adaptations

- **State Management:** NNX uses explicit state. Optimizer updates are now more functional (`optimizer.update(model, grads)`), making data flow clearer and easier to debug.
- **Error Messages:** JIT compilation can sometimes obscure error origins. Using `jax_debug_nans`, Chex, or `jax.disable_jit` helps pinpoint issues more accurately.

Recommended Debugging Workflow

A Systematic Approach to JAX/NNX Debugging

Static Checks First:

- Use `nnx.display()` to verify model/optimizer structure and initial state.
- Add Chex static assertions (`assert_shape`, `assert_type`) liberally.
- Manually check input/output shapes outside JIT.

A Systematic Approach to JAX/NNX Debugging

Runtime Issues within JIT:

- NaN/Inf? Use `chex.assert_tree_all_finite(model) + @chex.chexify`, or enable `jax_debug_nans`.
- Inspect values: Use `jax.debug.print`.
- Interactive step-through: Use `jax.debug.breakpoint`.

A Systematic Approach to JAX/NNX Debugging

Complex Issues / Need Standard Tools:

- Temporarily use with `jax.disable_jit()`: around the problematic code.
- Use standard `print()` and `pdb/breakpoint()`, or IDE debugger.

A Systematic Approach to JAX/NNX Debugging

Performance Issues:

- Suspect re-compilation? Use `chex.assert_max_traces`.
- Bottlenecks? Use `jax.profiler` (+ TensorBoard).

A Systematic Approach to JAX/NNX Debugging

Monitor Continuously:

- Integrate TensorBoard early for logging metrics and visualization.

Conclusion

- Debugging JAX/NNX requires adapting from PyTorch's eager-first model due to JIT.
- Master the JAX toolkit: `jax.debug.print`, `jax.debug.breakpoint`.
- Know when to use the "escape hatch": `jax.disable_jit`.
- Leverage Flax NNX's inspectability: `nnx.display()`.
- Build robustness with Chex assertions (static, value + `chex.chexify`).
- Monitor effectively with TensorBoard.
- Choose the right tool for the specific debugging task.

Learning Resources

Code Exercises, Quick References, and Slides

- <https://goo.gle/learning-jax>



Community and Docs

Community:

- <https://goo.gle/jax-community>

Docs

- JAX AI Stack: <https://jaxstack.ai>
- JAX: <https://jax.dev>
- Flax NNX: <https://flax.readthedocs.io>