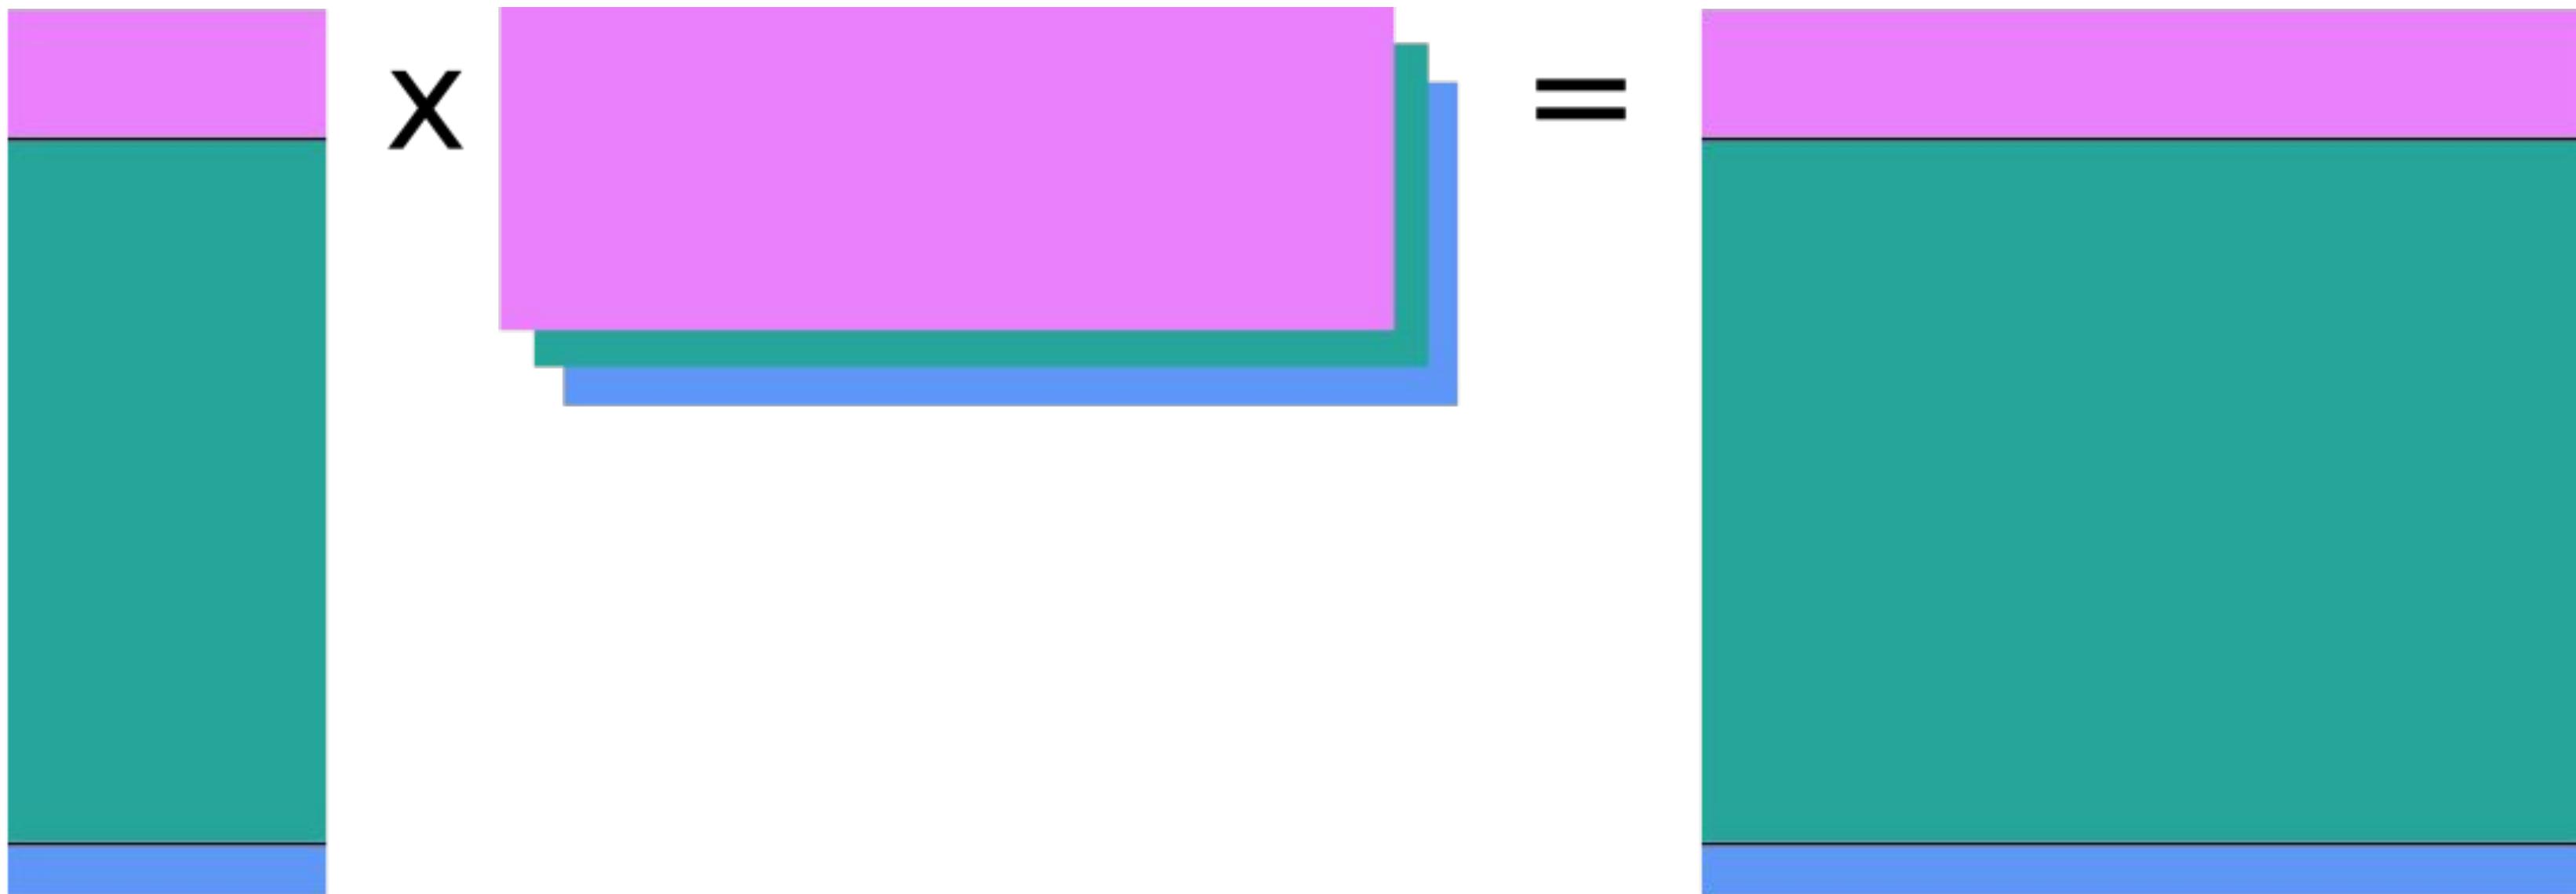


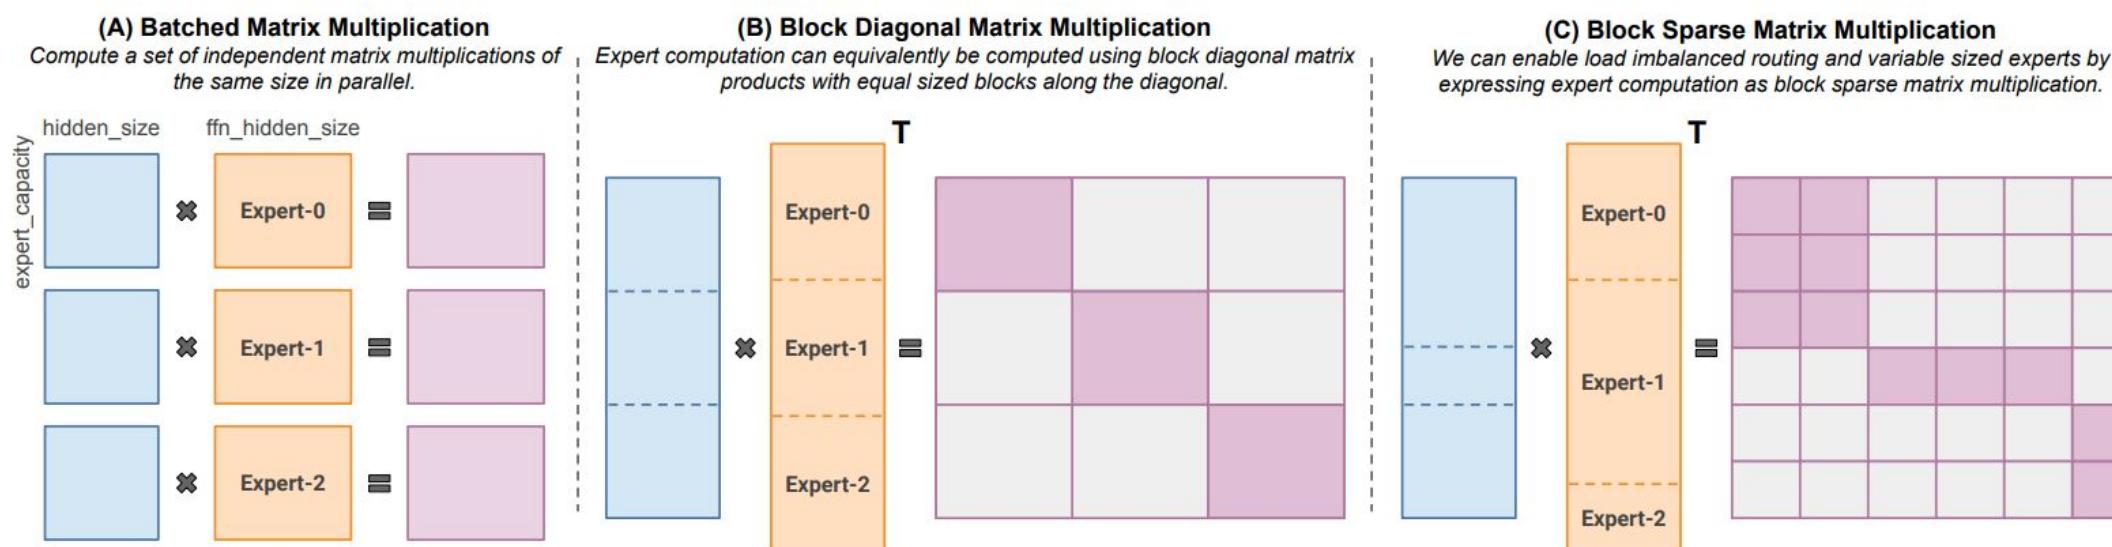
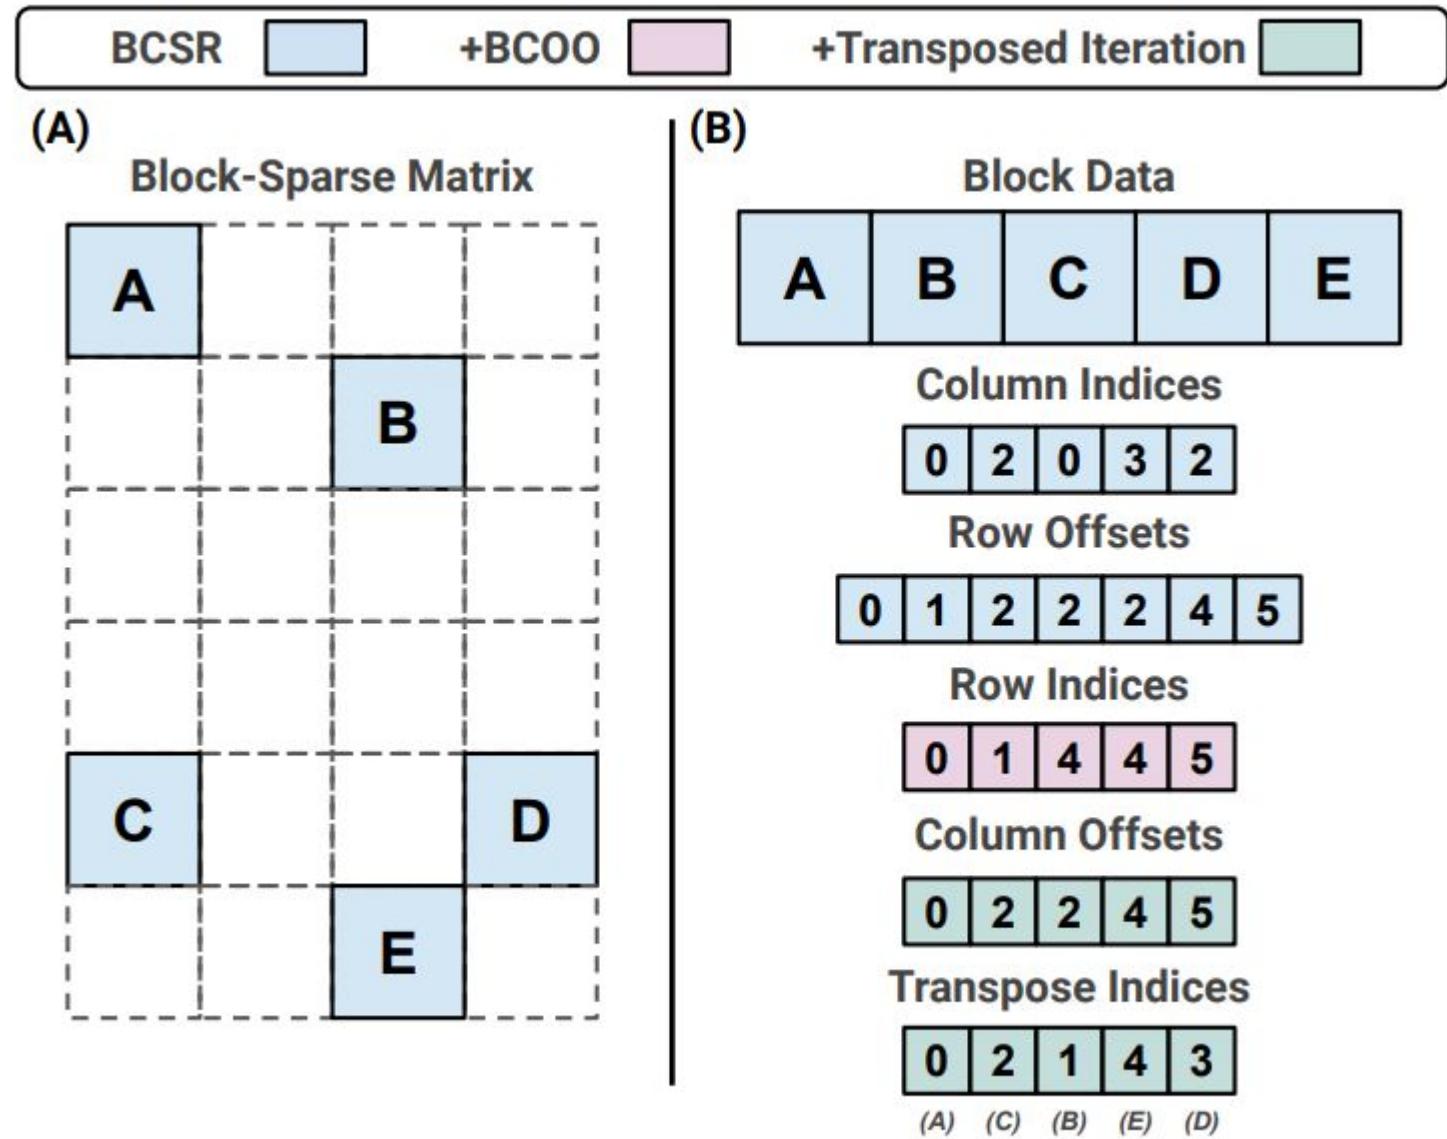
ragged dot



Intro

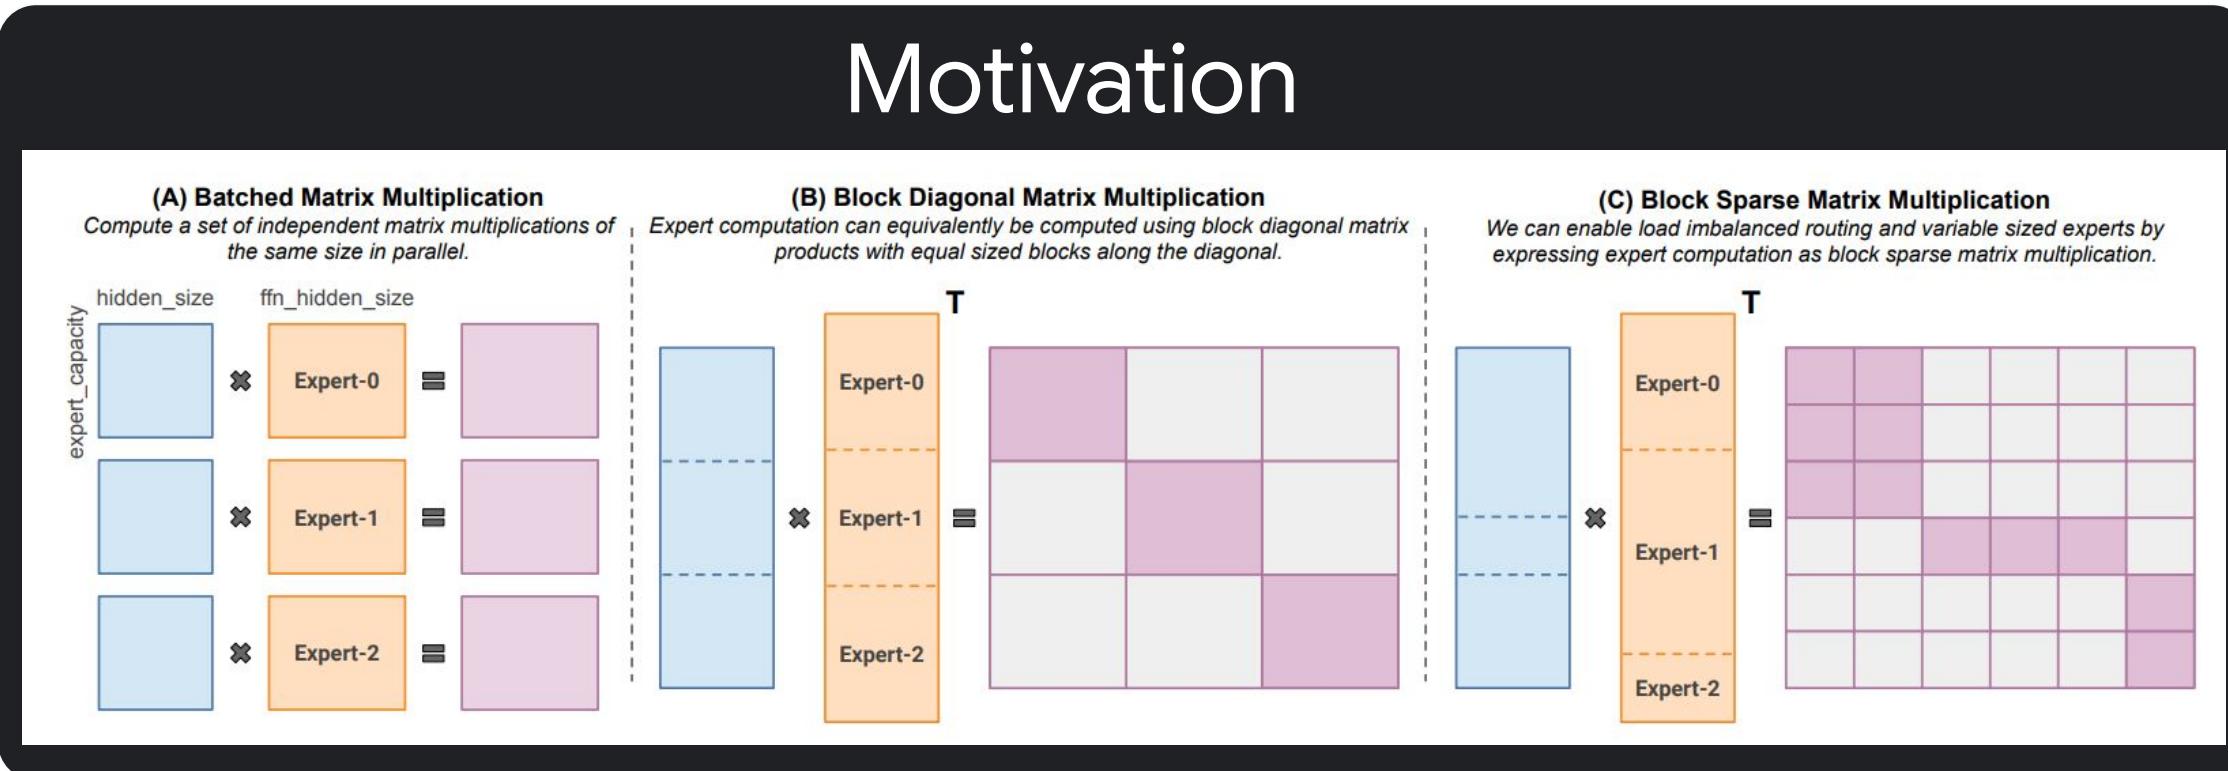
History

- *MegaBlocks: Efficient Sparse Training with Mixture-of-Experts*, Trevor Gale, Deepak Narayanan, Cliff Young, Matei Zaharia
- traditional sparsity represented using CSC or CSR (on GPU)
- modern accelerators deal badly with sparsity
- **Block-CSR - BCSR**



History

Motivation



- some possible approaches:
 - pad blocks to the maximum size (memory inefficient)
 - loop over rows individually (accelerator unfriendly)
- decent approaches on GPU
 - use NVIDIA libraries, dispatch **each expert on CPU**
 - not bad with a good compiler and/or CUDA-graphs

Base Routine

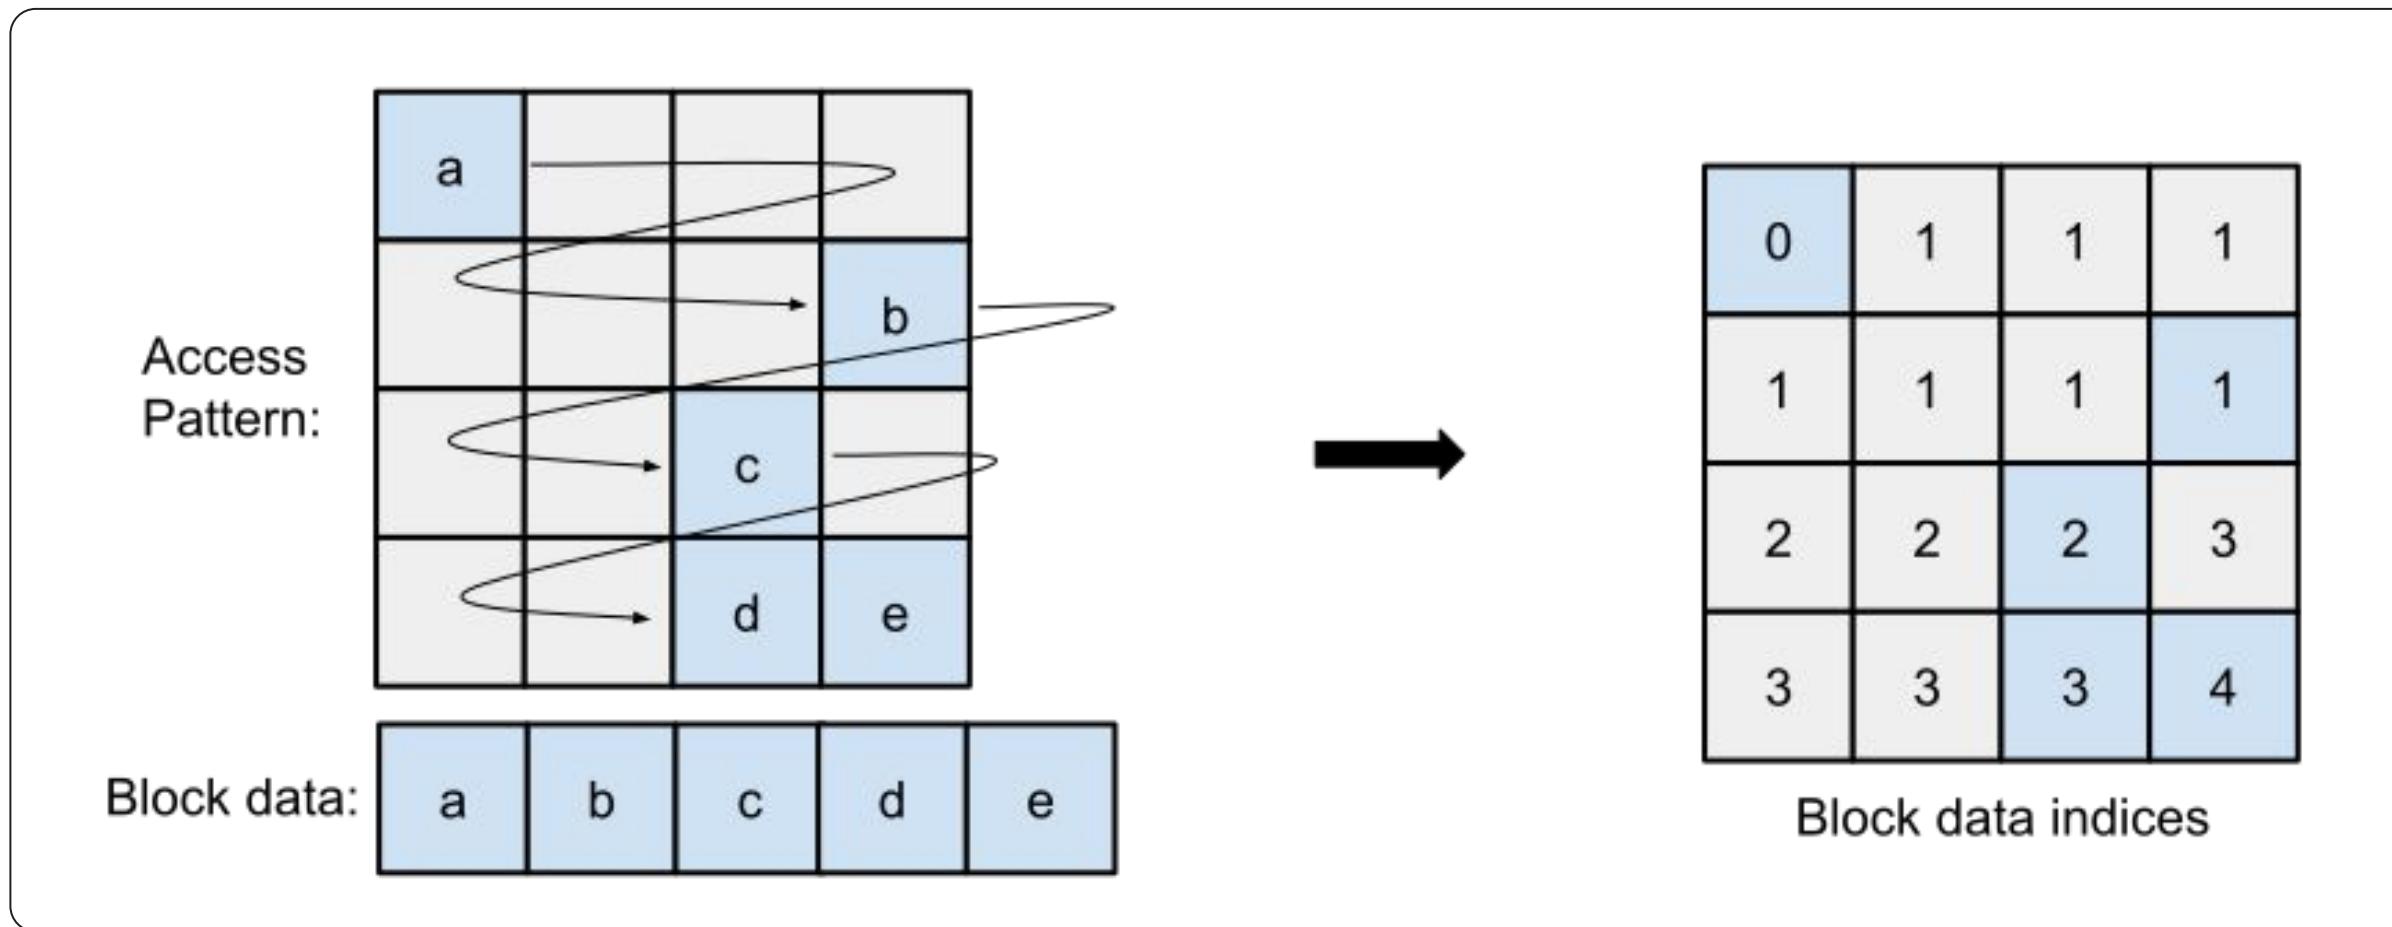
$$\begin{array}{|c|c|c|c|} \hline a & & & \\ \hline & & & \\ \hline & & c & \\ \hline & & d & e \\ \hline \end{array} @ \begin{array}{|c|c|c|c|} \hline w & & & \\ \hline x & & & \\ \hline y & & & \\ \hline z & & & \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline aw & & & \\ \hline bz & & & \\ \hline cy & & & \\ \hline dy+ez & & & \\ \hline \end{array}$$

Fig: Block-sparse matrix multiplication

docs.jax.dev/en/latest/pallas/tpu/sparse.html#example-sparse-dense-matrix-multiplication

Sparse computations in pallas

pltpu.PrefetchScalarGridSpec



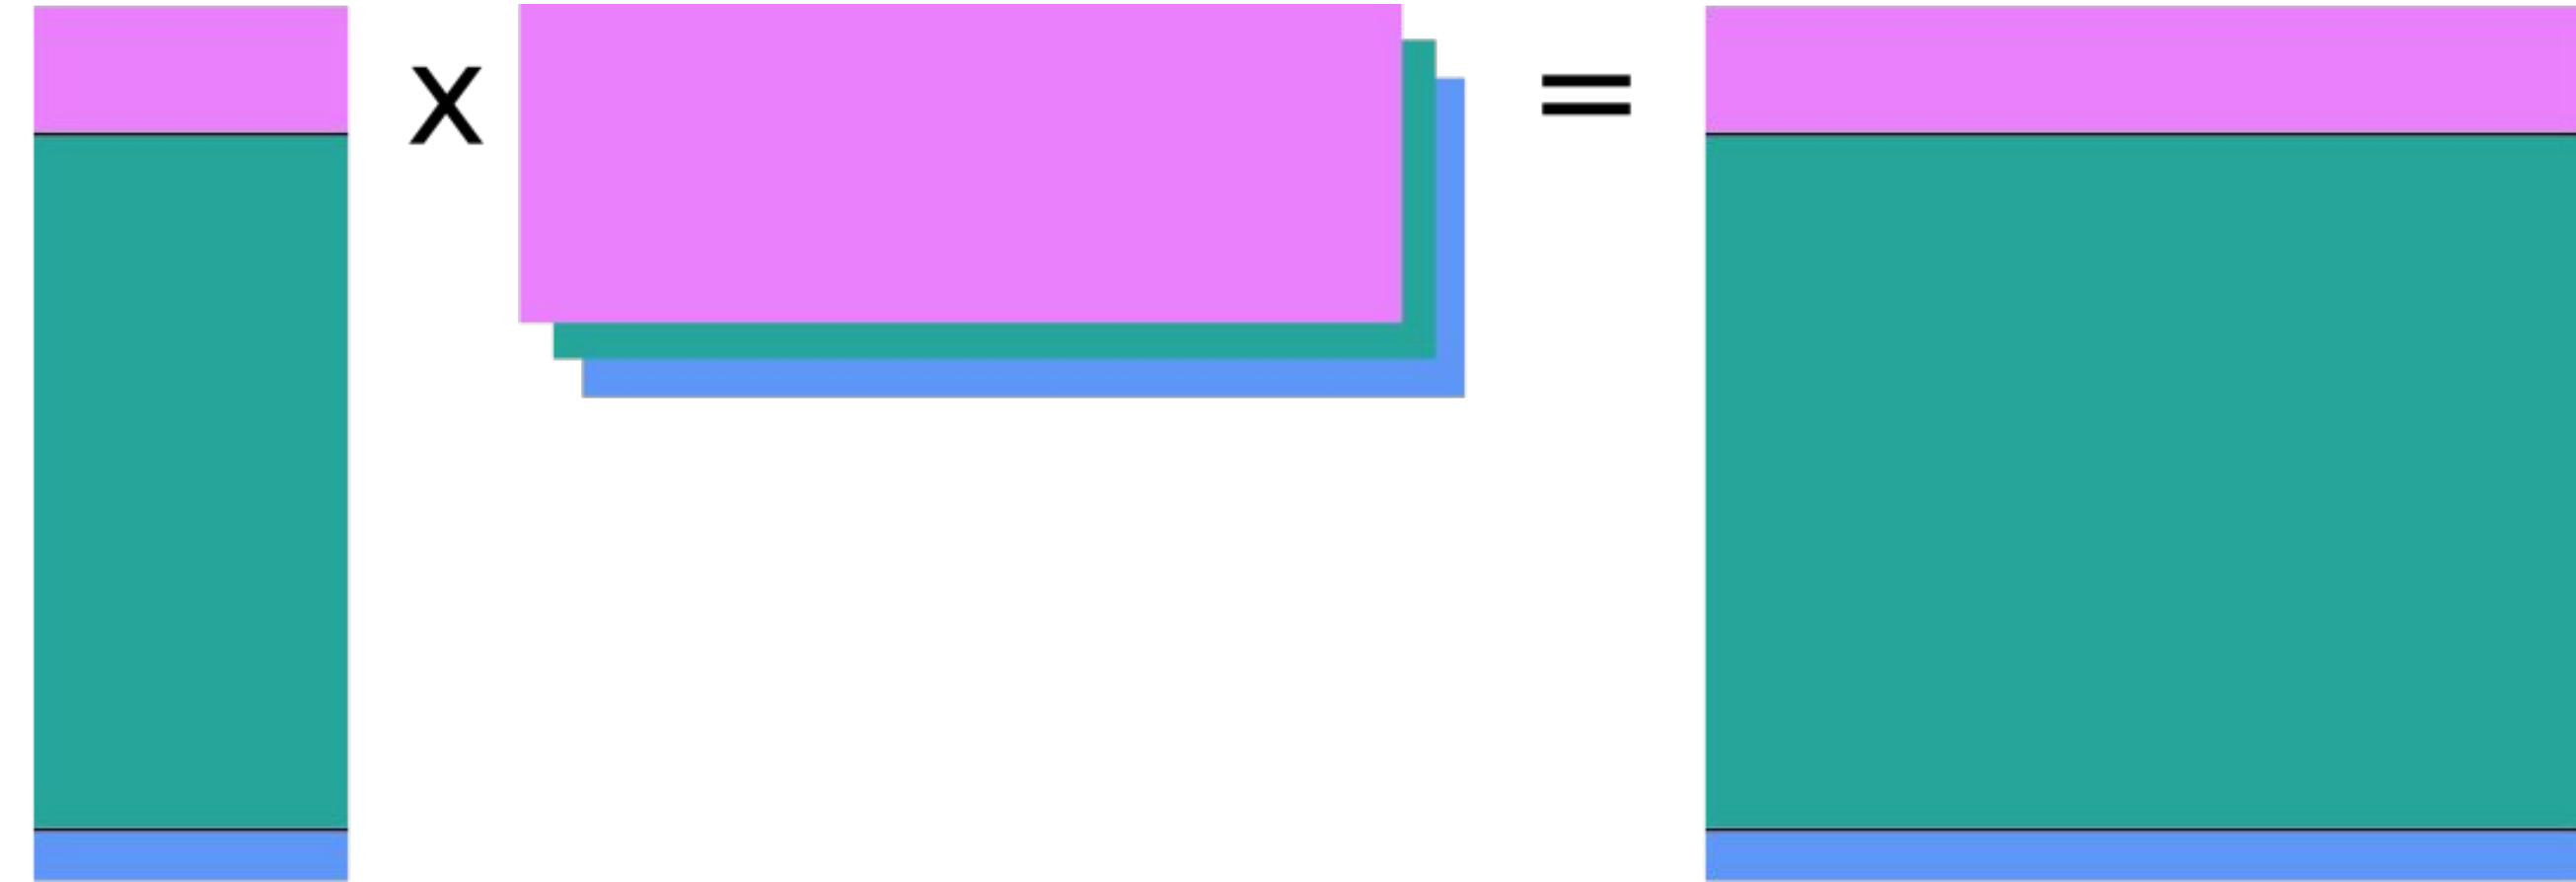
- pltpu.PrefetchScalarGridSpec is super useful
- generally pre-compute a metadata lookup table, then just iterate over entries
- example: `pl.BlockSpec(..., lambda i, j, scalar1_ref, scalar2_ref: (scalar1_ref[i], scalar2_ref[j]))`
- can re-implement the prefetch scalar grid using existing pallas APIs:

https://github.com/openxla/tokamax/blob/main/tokamax/_src/mosaic_tpu.py#L126

- “2, then 2” uses no BW
- read in the next block ASAP
- small penalty for visiting block at all
- can skip computation inside kernel

Modern “Megablox”

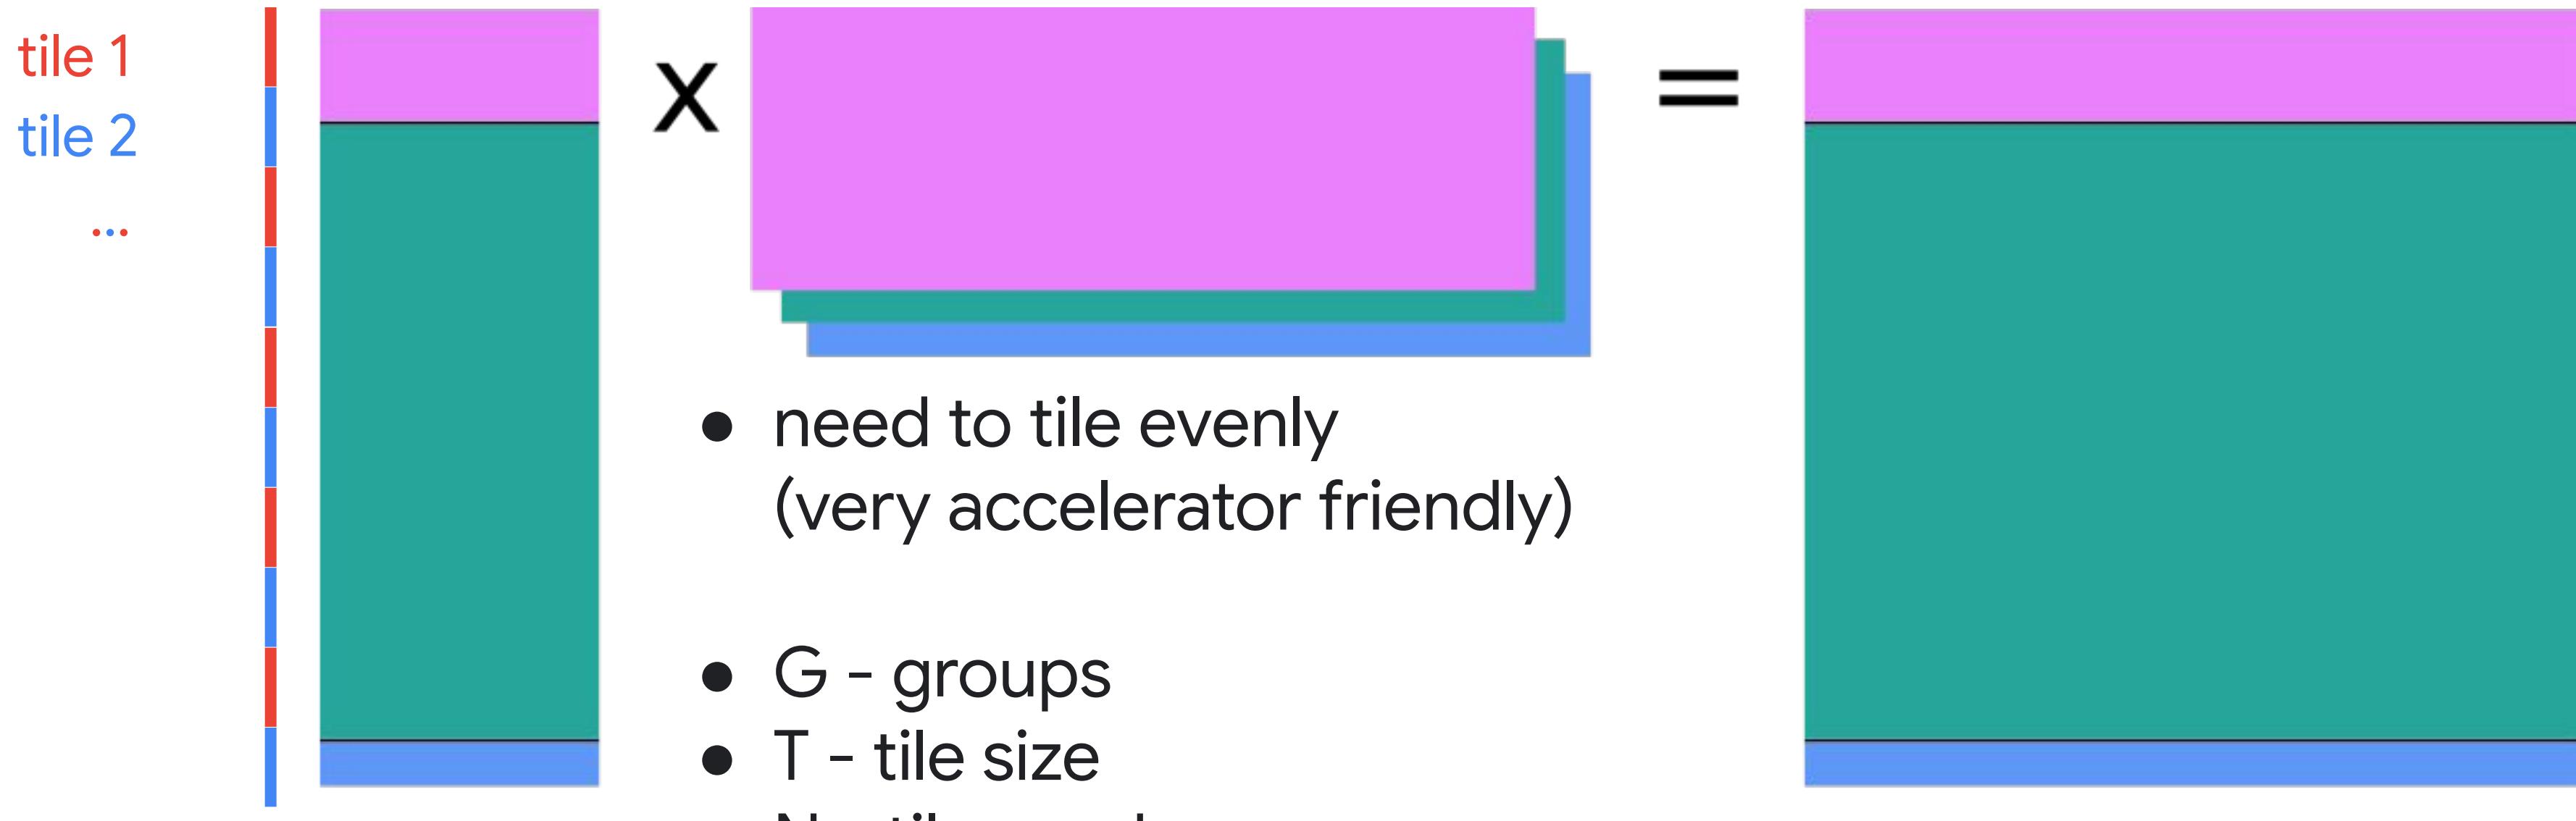
Megablox



- “ragged” representation usually means “pack stuff along one axis”
- the most dominant representation these days
- contiguous groups of rows belong to the (ordered) matrix in the stack
- LHS is 2D: $[m, k]$
- RHS is 3D: [group, k, n]
- OUT is 2D: $[m, n]$
- (theoretical) FLOPS are the same as for a simple matmul: $2 \cdot m \cdot k \cdot n$

Megablox

Accelerator Tiling & Tile Revisiting



- probability of having to “revisit” a tile: $P(\text{revisit}) = 1 - 1 / T$
- number of tiles revisited: $G \cdot P(\text{revisit}) = G \cdot (1 - 1 / T)$
- efficiency is tile number / actual tiles visited: $N / (N + G \cdot (1 - 1 / T))$

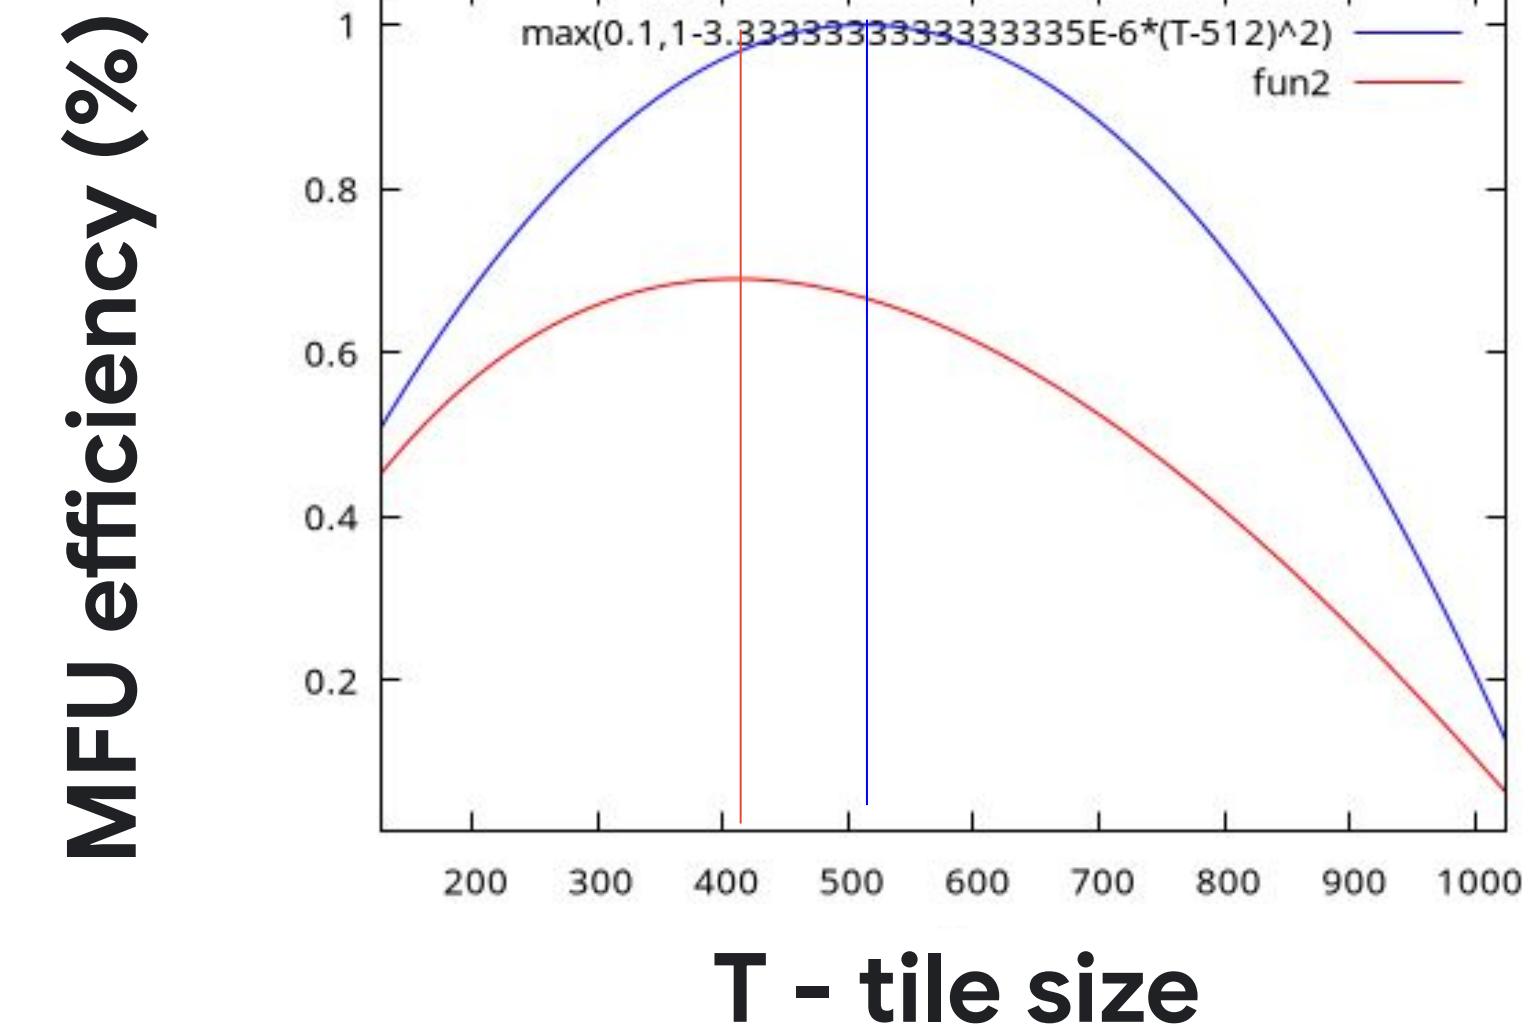
Megablox Accelerator Efficiency

- G - groups
- T - tile size
- N - tile number

- probability of having to “revisit” a tile: $P(\text{revisit}) = 1 - 1 / T$
- number of tiles revisited: $G \cdot P(\text{revisit}) = G \cdot (1 - 1 / T)$
- efficiency is tile number / actual tiles visited: $N / (N + G \cdot (1 - 1 / T))$

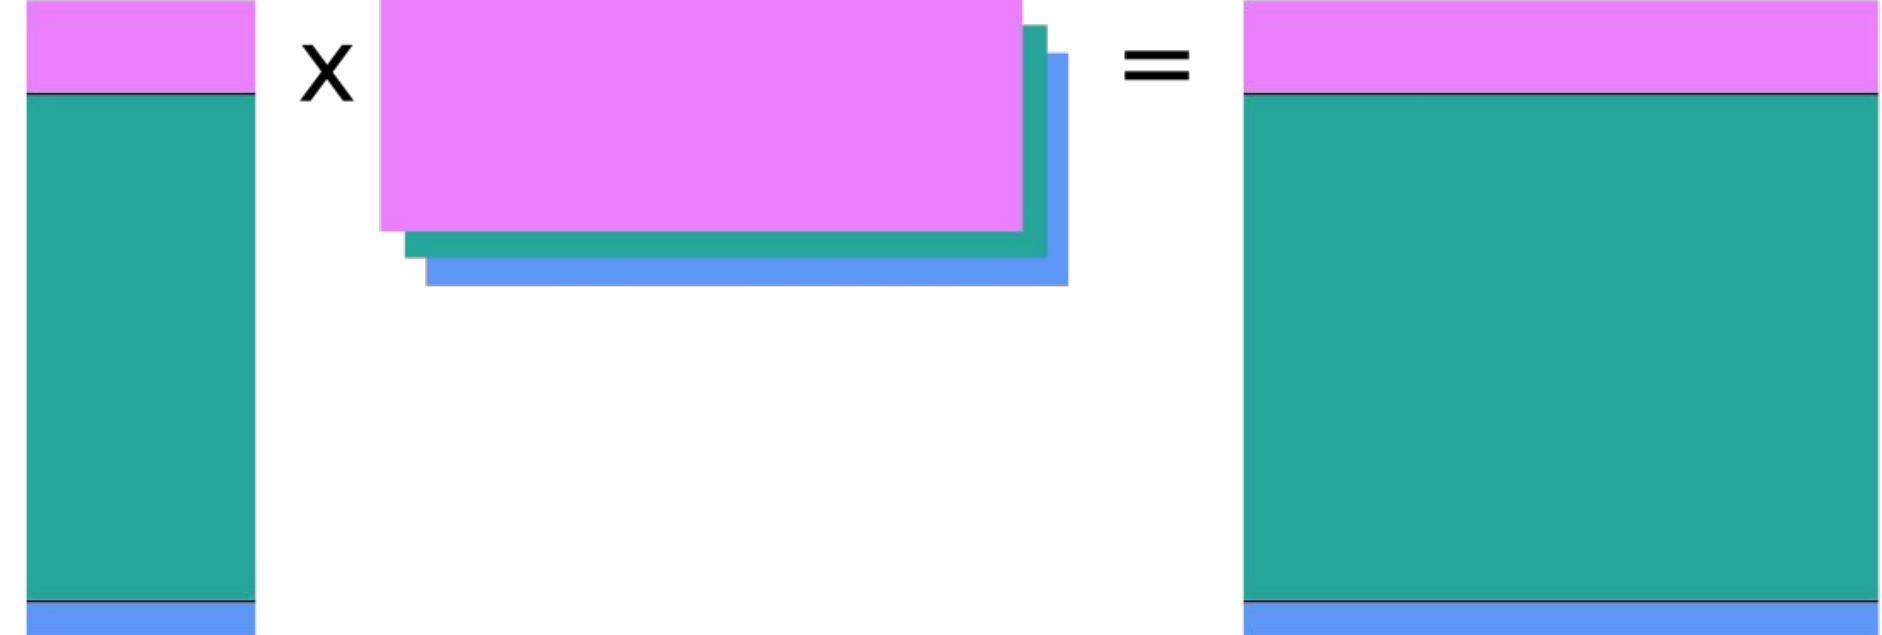
megablox efficiency model

- assume optimal matmul tile size model
- smaller tile sizes: more block revisiting
- overall lower MFU
- optimal megablox tile size skews smaller



Megablox Derivatives

- LHS: [m, k]
- RHS: [g, k, n]
- OUT: [m, n]



d(LHS)

- needs to be [m, k]
- `ragged_dot(`
 `d(OUT),`
 `RHST` # transpose RHS
 `)`
- FLOPS: $2 \cdot m \cdot k \cdot n$ (still)
- reduction along the n-axis

d(RHS)

- needs to be [g, k, n]
- “transposed”_`ragged_dot(`
 `LHST,`
 `d(OUT)`
 `)`
- FLOPS: $2 \cdot m \cdot k \cdot n$ (still)
- lots of outer products
- reduction along: **masked** m-axis

Megablox

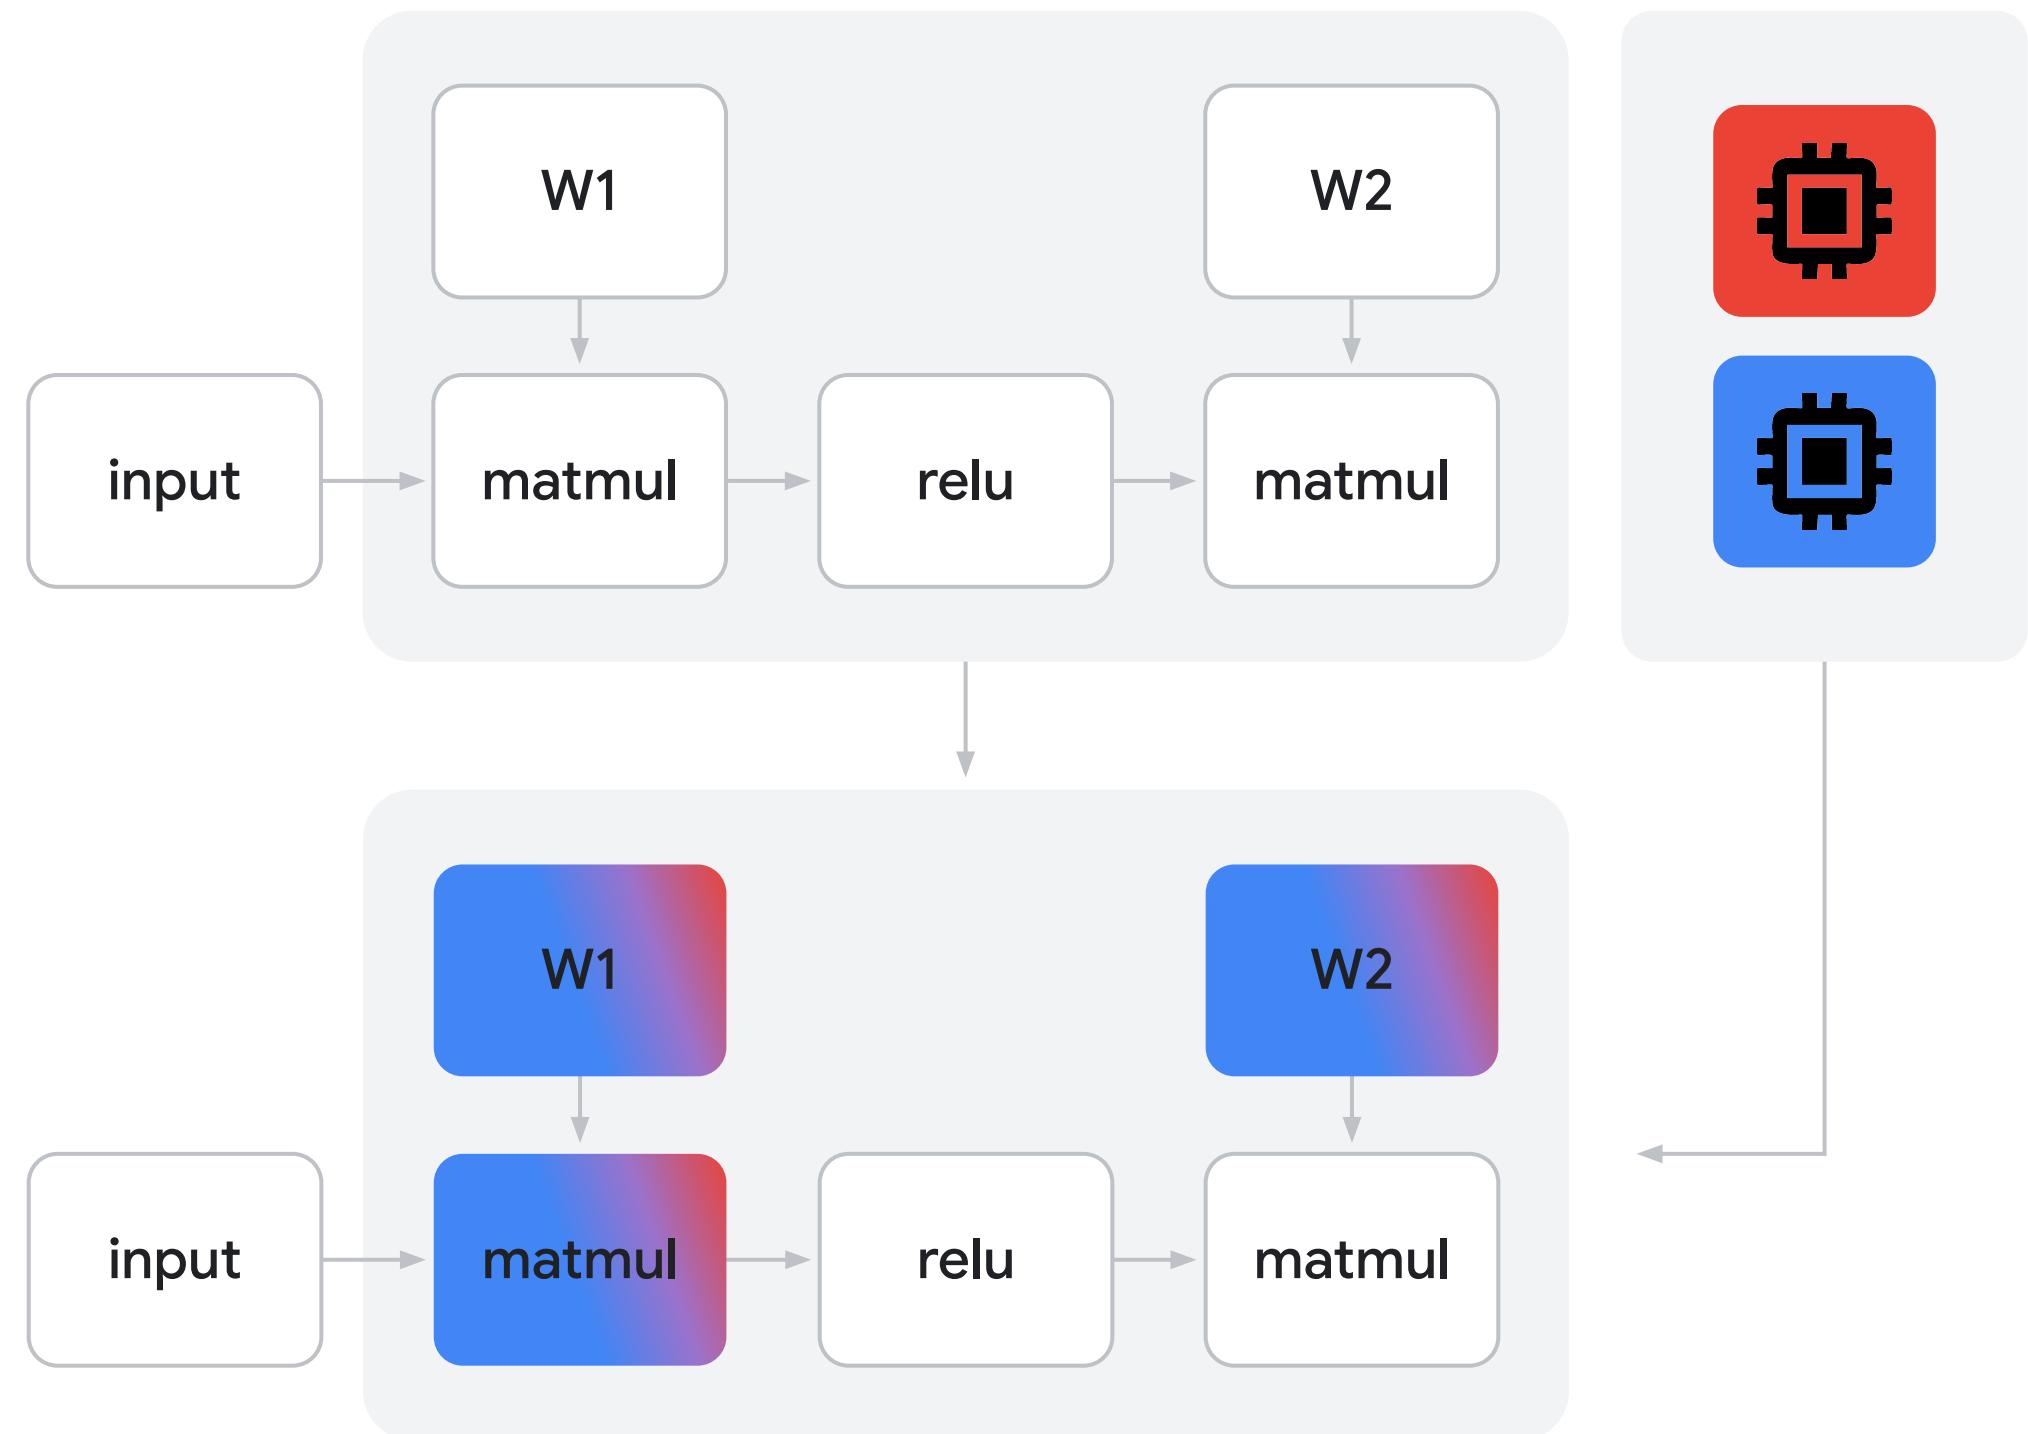
Actual metadata computation

- don't visit & skip empty blocks
 - this costs kernel dispatch latency
- compute tile and expert map flat
 - sometimes tile increase
 - sometimes expert
 - just a flat lookup table
- how to do this efficiently on TPU?
 - no loops, must use e.g., cumsum
 - computing metadata blocks TC

```
1 def make_group_metadata(
2     *,
3     group_sizes: jnp.ndarray,
4     m: int,
5     tm: int,
6     start_group: jnp.ndarray,
7     num_nonzero_groups: int,
8     visit_empty_groups: bool = True,
9 ) -> GroupMetadata:
10     """Create the metadata needed for grouped matmul computation.
11
12     Args:
13         group_sizes: A 1d, jnp.ndarray with shape [num_groups] and jnp.int32 dtype.
14         m: The number of rows in lhs.
15         tm: The m-dimension tile size being used.
16         start_group: The group in group sizes to start computing from. This is
17             particularly useful for when rhs num groups is sharded.
18         num_nonzero_groups: Number of groups in group sizes to compute on. Useful in
19             combination with group_offset.
20         visit_empty_groups: If True, do not squeeze tiles for empty groups out of
21             the metadata. This is necessary for tgmm, where we at least need to zero
22             the output for each group.
23
24     Returns:
25         tuple of:
26             group_offsets: A 1d, jnp.ndarray with shape [num_groups+1] and jnp.int32
27                 dtype. group_offsets[i] indicates the row at which group [i] starts in
28                 the lhs matrix and group_offsets[i-1] = m.
29             group_ids: A 1d, jnp.ndarray with shape [m_tiles + num_groups] and
30                 jnp.int32 dtype. group_ids[i] indicates which group grid index 'i' will
31                 work on.
32             m_tile_ids: A 1d, jnp.ndarray with shape [m_tiles + num_groups] and
33                 jnp.int32. m_tile_ids[i] indicates which m-dimension tile grid index 'i'
34                 will work on.
35             num_tiles: The number of m-dimension tiles to execute.
36
37     num_groups = group_sizes.shape[0]
38     end_group = start_group + num_nonzero_groups - 1
39
40     # Calculate the offset of each group, starting at zero. This metadata is
41     # similar to row offsets in a CSR matrix. The following properties hold:
42     #
43     # group_offsets.shape = [num_groups + 1]
44     # group_offsets[0] = 0
45     # group_offsets[num_groups] = m
46     #
47     # The row at which group 'i' starts is group_offsets[i].
48     group_ends = jnp.cumsum(group_sizes)
49     group_offsets = jnp.concatenate([jnp.zeros(1, dtype=jnp.int32), group_ends])
50
51     # Assign a group id to each grid index.
52     #
53     # If a group starts somewhere other than the start of a tile or ends somewhere
54     # other than the end of a tile we need to compute that full tile. Calculate
55     # the number of tiles for each group by rounding their end up to the nearest
56     # 'tm' and their start down to the nearest 'tm'.
57
58     # (1) Round the group_ends up to the nearest multiple of 'tm'.
59     #
60     # NOTE: This does not change group_offsets[num_groups], which is m
61     # (because we enforce m is divisible by tm).
62     rounded_group_ends = ((group_ends + tm - 1) // tm * tm).astype(jnp.int32)
63
64     # (2) Round the group_starts down to the nearest multiple of 'tm'.
65     group_starts = jnp.concatenate(
66         [jnp.zeros(1, dtype=jnp.int32), group_ends[:-1]])
67
68     rounded_group_starts = group_starts // tm * tm
69
70     # (3) Calculate the number of rows in each group.
71     #
72     # NOTE: Handle zero-sized groups as a special case. If the start for a
73     # zero-sized group is not divisible by 'tm' its start will be rounded down and
74     # its end will be rounded up such that its size will become 1 tile here.
75     rounded_group_sizes = rounded_group_ends - rounded_group_starts
76     rounded_group_sizes = jnp.where(group_sizes == 0, 0, rounded_group_sizes)
77
78     # (4) Convert the group sizes from units of rows to unit of 'tm' sized tiles.
79     #
80     # An m-dimension tile is 'owned' by group 'i' if the first row of the tile
81     # belongs to group 'i'. In addition to owned tiles, each group can have 0 or 1
82     # initial partial tiles if it's first row does not occur in the first row of a
83     # tile. The '0-th' group never has a partial tile because it always starts at
84     # the 0-th row.
85
86     # If no group has a partial tile, the total number of tiles is equal to
87     # 'm // tm'. If every group has a partial except the 0-th group, the total
88     # number of tiles is equal to 'm // tm + num_groups - 1'. Thus we know that
89     #
90     # tiles_m <= group_tiles.sum() <= tiles_m + num_groups - 1
91
92     # Where tiles_m = m // tm.
93     #
94     # NOTE: All group sizes are divisible by 'tm' because of the rounding in steps
95     # (1) and (2) so this division is exact.
96     group_tiles = rounded_group_sizes // tm
97
98     if visit_empty_groups:
99         # Insert one tile for empty groups.
100        group_tiles = jnp.where(group_sizes == 0, 1, group_tiles)
101
102    # Create the group ids for each grid index based on the tile counts for each
103    # group.
104
105    # NOTE: This repeat(...) will pad group_ids with the final group id if
106    # group_tiles.sum() < tiles_m + num_groups - 1. The kernel grid will be sized
107    # such that we only execute the necessary number of tiles.
108    tiles_m = _calculate_num_tiles(m, tm)
109    group_ids = jnp.repeat(
110        jnp.arange(num_groups, dtype=jnp.int32),
111        group_tiles,
112        total_repeat_length=tiles_m + num_groups - 1,
113    )
114
115    # Assign an m-dimension tile id to each grid index.
116    #
117    # NOTE: Output tiles can only be re-visited consecutively. The following
118    # procedure guarantees that m-dimension tile indices respect this.
119
120    # (1) Calculate how many times each m-dimension tile will be visited.
121    #
122    # Each tile is guaranteed to be visited once by the group that owns the tile.
123    # The remaining possible visits occur when a group starts inside of a tile at
124    # a position other than the first row. We can calculate which m-dimension tile
125    # each group starts in by floor-dividing its offset with 'tm' and then count
126    # tile visits with a histogram.
127    #
128    # To avoid double counting tile visits from the group that owns the tile,
129    # filter these out by assigning their tile id to 'tile_m' (one beyond the max)
130    # such that they're ignored by the subsequent histogram. Also filter out any
131    # group which is empty.
132    #
133    # TODO(tgale): Invert the 'partial_tile_mask' predicates to be more clear.
134    partial_tile_mask = jnp.logical_or(
135        (group_offsets[:-1] % tm) == 0, group_sizes == 0
136    )
137
138    # Explicitly enable tiles for zero sized groups, if specified. This covers
139    # zero sized groups that start on a tile-aligned row and those that do not.
140    if visit_empty_groups:
141        partial_tile_mask = jnp.where(group_sizes == 0, 0, partial_tile_mask)
142
143    partial_tile_ids = jnp.where(
144        partial_tile_mask, tiles_m, group_offsets[:-1] // tm
145    )
146
147    tile_visits = (
148        jnp.histogram(partial_tile_ids, bins=tiles_m, range=(0, tiles_m - 1))[0]
149        + 1
150    )
151
152    # Create the m-dimension tile ids for each grid index based on the visit
153    # counts for each tile.
154    m_tile_ids = jnp.repeat(
155        jnp.arange(tiles_m, dtype=jnp.int32),
156        tile_visits.astype(jnp.int32),
157        total_repeat_length=tiles_m + num_groups - 1,
158    )
159
160    # Account for sharding.
161    #
162    # Find the start of the groups owned by our shard and shift the group_ids and
163    # m_tile_ids s.t. the metadata for our tiles are at the front of the arrays.
164    #
165    # TODO(tgale): Move this offset into the kernel to avoid these rolls.
166    first_tile_in_shard = (group_ids < start_group).sum()
167    group_ids = jnp.roll(group_ids, shift=-first_tile_in_shard, axis=0)
168    m_tile_ids = jnp.roll(m_tile_ids, shift=-first_tile_in_shard, axis=0)
169
170    # Calculate the number of tiles we need to compute for our shard.
171    #
172    # Remove tile visits that belong to a group not in our shard.
173    iota = jnp.arange(num_groups, dtype=jnp.int32)
174    active_group_mask = jnp.logical_and(iota <= end_group, iota >= start_group)
175    group_tiles = jnp.where(active_group_mask, group_tiles, 0)
176    num_tiles = group_tiles.sum()
177    return (group_offsets, group_ids, m_tile_ids), num_tiles
178
```

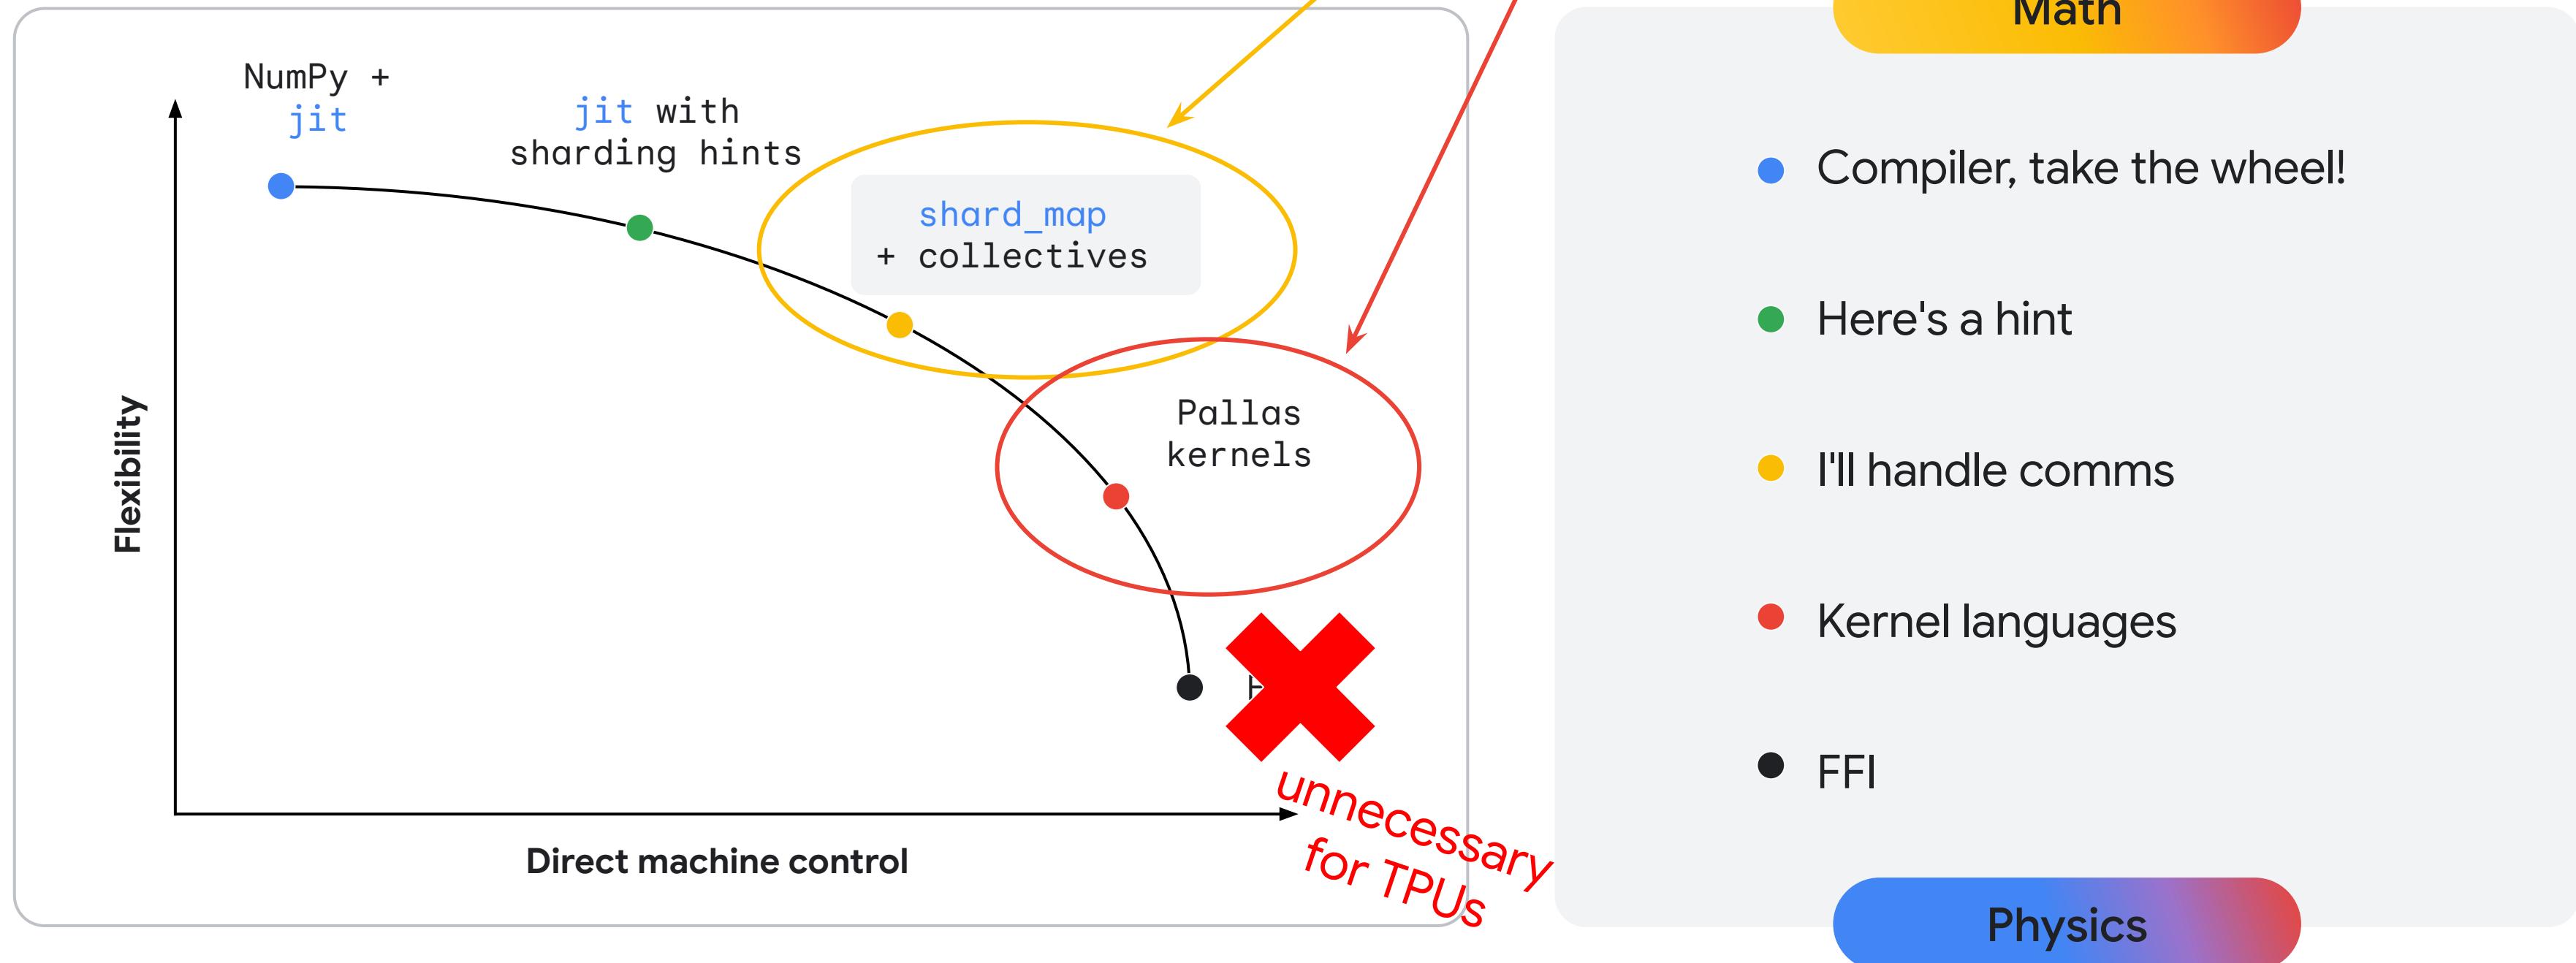
The MoE Layer in JAX

Sharded MLP (or MoE)



JAX: The escape hatch hierarchy

Flexibility vs. control



MoE Prep

helper lambdas

token routing
(just computing the index map)

in specs and out specs for shard_map

constants for the MoE closure

```
1 def moe_block(x: jax.Array, layer: MoELayer, cfg: Config):
2     assert x.ndim == 3
3     l2p = lambda *axes: logical_to_physical(axes, cfg.rules)
4     _psc = lambda z, spec: reshard(z, P(*spec))
5     psc = lambda z, spec: _qpsc(z, spec) if is_type(z, QuantArray) else _psc(z, spec)
6
7     # we're decoding or device count does not divide total token count
8     replicated_routing = x.shape[-2] == 1 or (x.shape[-2] * x.shape[-3]) % jax.device_count() != 0
9     topk_weights, topk_idx = _route_tokens_to_moe_experts(x, layer.w_router, replicated_routing, cfg)
10    tensor_axname, expert_axname = l2p("moe_e_tp")[0], l2p("moe_e_ep")[0]
11
12    x_spec = l2p("batch", "sequence", None)
13    topk_weights_spec, topk_idx_spec = l2p("batch", "sequence", None), l2p("batch", "sequence", None)
14    out_spec = l2p("batch", "sequence", None)
15
16    we_gate_spec = l2p("moe_e_ep", None, "moe_e_tp")
17    we_up_spec = l2p("moe_e_ep", None, "moe_e_tp")
18    we_down_spec = l2p("moe_e_ep", "moe_e_tp", None)
19    we_gate = psc(layer.we_gate, we_gate_spec)
20    we_up = psc(layer.we_up, we_up_spec)
21    we_down = psc(layer.we_down, we_down_spec)
22
23    in_specs = (x_spec, we_gate_spec, we_up_spec, we_down_spec, topk_weights_spec, topk_idx_spec)
24
25    is_embedding_sharded = l2p("act_embed")[0] is not None
26    if is_embedding_sharded: # activations are sharded
27        out_spec = P(*out_spec[:-1] + (tensor_axname,)) # override last axis name
28
29    expert_count = cfg.mesh.axis_sizes[cfg.mesh.axis_names.index(expert_axname)] if expert_axname is not None else 1
30    tensor_count = cfg.mesh.axis_sizes[cfg.mesh.axis_names.index(tensor_axname)] if tensor_axname is not None else 1
31    assert cfg.moe_numExperts % expert_count == 0
32    expert_size = cfg.moe_numExperts // expert_count
```

MoE Compute

sort indices

gather tokens
(potentially increase their count)

group sizes via bincount

up and gate projection
(first tensor-parallel stage)

down projection
(second tensor-parallel stage)

weight the tokens

all-gather tokens
and reduce across experts

```
38
39 @partial(shard_map, mesh=cfg.mesh, in_specs=in_specs, out_specs=out_spec, check_rep=False)
40 def _expert_fn(x, we_gate, we_up, we_down, topk_weights, topk_idx):
41     (b, s, d), e = x.shape, cfg.moe_experts_per_tok
42     expert_idx = jax.lax.axis_index(expert_axname) if expert_axname is not None else 0
43     tensor_idx = jax.lax.axis_index(tensor_axname) if tensor_axname is not None else 0
44     del tensor_idx
45     topk_idx_ = topk_idx.reshape(-1)
46     valid_group_mask_ = (topk_idx_ >= expert_size * expert_idx) & (topk_idx_ < expert_size * (expert_idx + 1))
47     expert_mapped_topk_idx_ = jnp.where(valid_group_mask_, topk_idx_ - expert_idx * expert_size, 2**30)
48
49     sort_idx_ = jnp.argsort(expert_mapped_topk_idx_, axis=-1) # [b * s * e]
50     isort_idx_ = jnp.argsort(sort_idx_)
51
52     topk_idx_sort_ = topk_idx_[sort_idx_] # [b * s * e]
53     expert_mapped_topk_idx_sort_ = expert_mapped_topk_idx_[sort_idx_]
54     valid_group_mask_sort_ = expert_mapped_topk_idx_sort_ < 2**30
55     expert_mapped_topk_idx_sort_ = jnp.where(expert_mapped_topk_idx_sort_ < 2**30, expert_mapped_topk_idx_sort_, 0)
56
57     # equivalent to:
58     #
59     # x_repeat_ = jnp.repeat(x.reshape((-1, x.shape[-1])), e, axis=0)
60     # x_repeat_sort_ = jnp.take_along_axis(x_repeat_, sort_idx_[:, None], axis=-2) # [b * s, d]
61     #
62     x_repeat_sort_ = jnp.take_along_axis(
63         x.reshape((-1, x.shape[-1])),
64         sort_idx_[:, None] // e,
65         axis=-2, # index trick to avoid jnp.repeat
66     ) # [b * s * e, d]
67
68     group_sizes = jnp.bincount(topk_idx_sort_, length=cfg.moe_numExperts)
69     group_sizes_shard = jax.lax.dynamic_slice_in_dim(group_sizes, expert_idx * expert_size, expert_size, 0)
70
71     with jax.named_scope("we_gate"):
72         ff_gate = _moe_gmm(x_repeat_sort_, we_gate, group_sizes_shard, expert_mapped_topk_idx_sort_, cfg)
73         ff_gate = jax.nn.silu(ff_gate)
74         ff_gate = jnp.where(valid_group_mask_sort_..., None, ff_gate, 0)
75     with jax.named_scope("we_up"):
76         ff_up = _moe_gmm(x_repeat_sort_, we_up, group_sizes_shard, expert_mapped_topk_idx_sort_, cfg)
77         ff_gate_up = jnp.where(valid_group_mask_sort_..., None, ff_gate * ff_up, 0)
78     with jax.named_scope("we_down"):
79         ff_out = _moe_gmm(ff_gate_up, we_down, group_sizes_shard, expert_mapped_topk_idx_sort_, cfg)
80         ff_out = jnp.where(valid_group_mask_sort_..., None, ff_out, 0) # expensive
81
82     ff_out = ff_out * topk_weights.reshape(-1)[sort_idx_[:, None]]
83
84     with jax.named_scope("unpermute"):
85         ff_out = jnp.take_along_axis(ff_out, isort_idx_..., None, axis=-2)
86     with jax.named_scope("expert_summing"):
87         ff_out_expert = jnp.sum(ff_out.reshape((b * s, e, d)), -2)
88         ff_out_expert = ff_out_expert.astype(cfg.dtype)
89
90     with jax.named_scope("experts_collective"):
91         # collectives
92         psum_axes = tensor_axname if expert_axname is None else (expert_axname, tensor_axname)
93         ff_out_expert = jax.lax.psum(ff_out_expert, psum_axes)
94         ff_out_expert = ff_out_expert.reshape((b, s, ff_out_expert.shape[-1]))
95     return ff_out_expert
96
97 with jax.named_scope("moe_routed_expert"):
98     x_ = psc(x, x_spec)
99     ff_out_expert = _expert_fn(x_, we_gate, we_up, we_down, topk_weights, topk_idx)[..., :x_.shape[-1]]
100 return psc(ff_out_expert, l2p("batch", "sequence", "act_embed"))
```

Quantization

Quantizing Matmuls

most hardware computes matmuls 2x as fast in lower precision (int8 / fp8)

full-channel quantization:

- compute-bound: $([m, k], [m, 1]) @ ([k, n], [1, n]) = ([m, k] @ [k, n]) * [m, 1] * [1, n]$
- HBM BW-bound:
 - option 1 (scale out): $[m, k] @ ([k, n], [1, n]) = ([m, k] @ [k, n]) * [1, n]$
 - option 1 (scale in): $[m, k] @ ([k, n], [k, 1]) = ([m, k] * [k, 1]) @ [k, n]$

what about subchannel quantization?

- similar, but each tile size along reduction dimension receives separate scale
- $([\text{tile}_m, \text{tile}_k] @ [\text{tile}_k, \text{tile}_n]) * [\text{tile}_m, 1] * [1, \text{tile}_n]$

Quantizing Ragged Dot

OUT: `ragged_dot(LHS, RHS)`

`d(LHS)`: `ragged_dot(d(OUT), RHST)`

`d(RHS)`: `transposed_ragged_dot(LHST, d(OUT))`

LHS: [m, k]
RHS: [g, k, n]
OUT: [m, n]

- quantizing very similar to matmul quantization strategies
 - full-channel quantization:
 - $([m, k], [m, 1]) @ ([k, n], [1, n]) = ([m, k] @ [k, n]) * [m, 1] * [1, n]$
-

- usually want single scale (or just a few scales, subchannel) along **reduction axis**
- reduction axis **changes** (matmuls have the same problem)

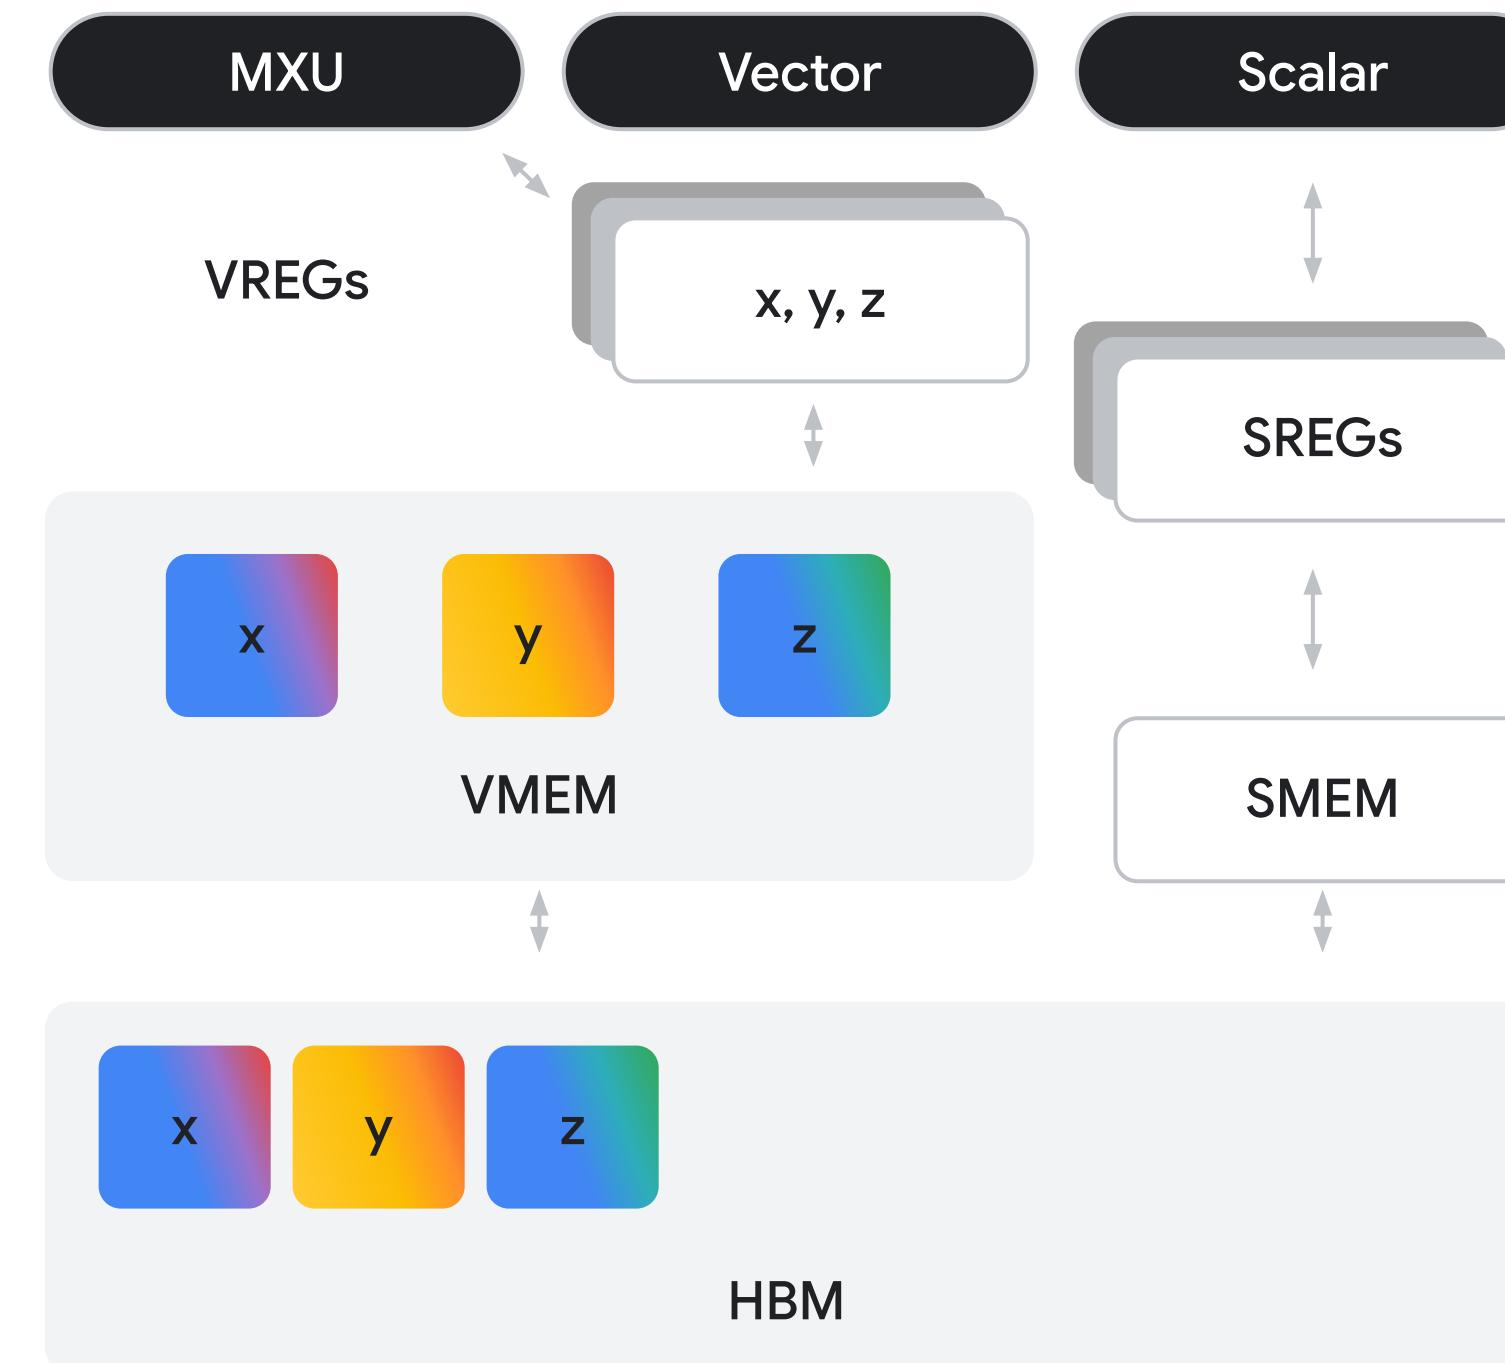
OUT: k-axis

`d(LHS)`: n-axis

`d(RHS)`: m-axis

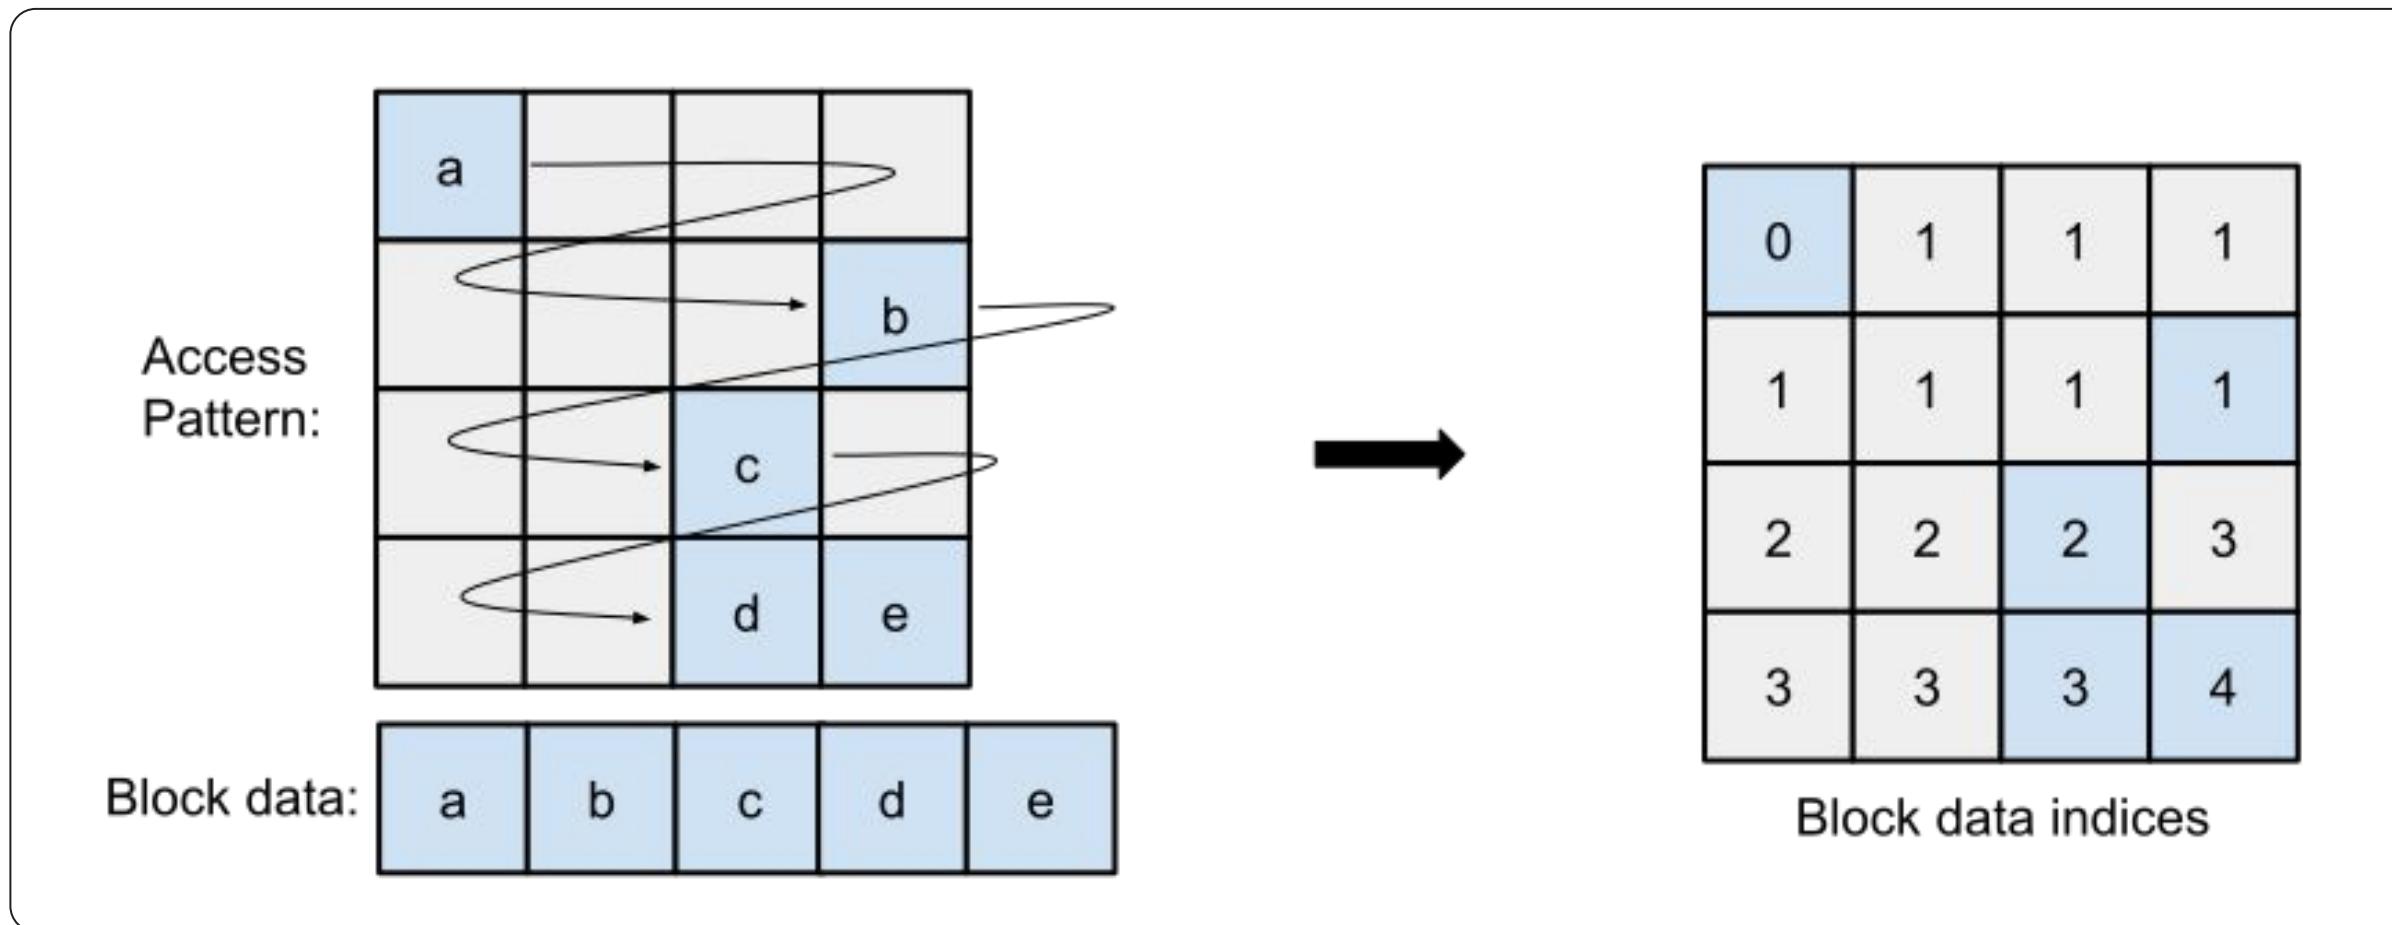
Matmuls & Megablox in Pallas

TPU Memory pipeline



Sparse computations in pallas

pltpu.PrefetchScalarGridSpec



- pltpu.PrefetchScalarGridSpec is super useful
- generally pre-compute a metadata lookup table, then just iterate over entries
- example: `pl.BlockSpec(..., lambda i, j, scalar1_ref, scalar2_ref: (scalar1_ref[i], scalar2_ref[j]))`
- can re-implement the prefetch scalar grid using existing pallas APIs:

https://github.com/openxla/tokamax/blob/main/tokamax/_src/mosaic_tpu.py#L126

- “2, then 2” uses no BW
- read in the next block ASAP
- small penalty for visiting block at all
- can skip computation inside kernel

Writing a Matmul on TPUs

- a super simple kernel
 - `pl.dot` to get correct “output” dtype
 - we’re issuing a hardware matmul
- 3D grid, no need for anything special
- **megablox** kernel is essentially this
 - apart from metadata of course
 - some row masking for tile overlap

```
def matmul_kernel(x_ref, y_ref, out_ref):
    out_ref[...] = pl.dot(x_ref[...], y_ref[...]).astype(out_ref.dtype)

def matmul(A, B, block_m, block_n, block_k):
    in_specs = [
        pl.BlockSpec((block_m, block_k), lambda i, j, k: (i, k)),
        pl.BlockSpec((block_k, block_n), lambda i, j, k: (k, j))
    ]
    out_specs = pl.BlockSpec((block_m, block_n), lambda i, j, k: (i, j))
    return pl.pallas_call(
        matmul_kernel,
        out_shape=jax.ShapeDtypeStruct((m, n), "bfloat16"),
        grid=(pl.cdiv(m, block_m), pl.cdiv(n, block_n), pl.cdiv(k, block_k)),
        in_specs=in_specs, out_specs=out_specs,
    )(A, B)
c = jax.jit(partial(matmul, block_m=2048, block_n=1024, block_k=1024))(A, B)
```

matmul: gist.github.com/rdyro/dd149ef5650f185bd96ec0666a23de9e
megablox: github.com/openxla/tokamax ... pallas_mosaic_tpu_kernel.py

Tuning is (essentially) necessary

tune-jax

```
import tune_jax
tune_jax.logger.setLevel("INFO") # print some info for sanity
tiles = [512, 1024, 2048, 4096] # any multiple of 128 will do
hyperparams = dict(
    block_m=tiles,
    block_k=tiles,
    block_n=tiles,
)
fn = tune_jax.tune(matmul, hyperparams=hyperparams)
_ = fn(A, B) # run to tune
print(tune_jax.tabulate(fn)) # print results nicely
```

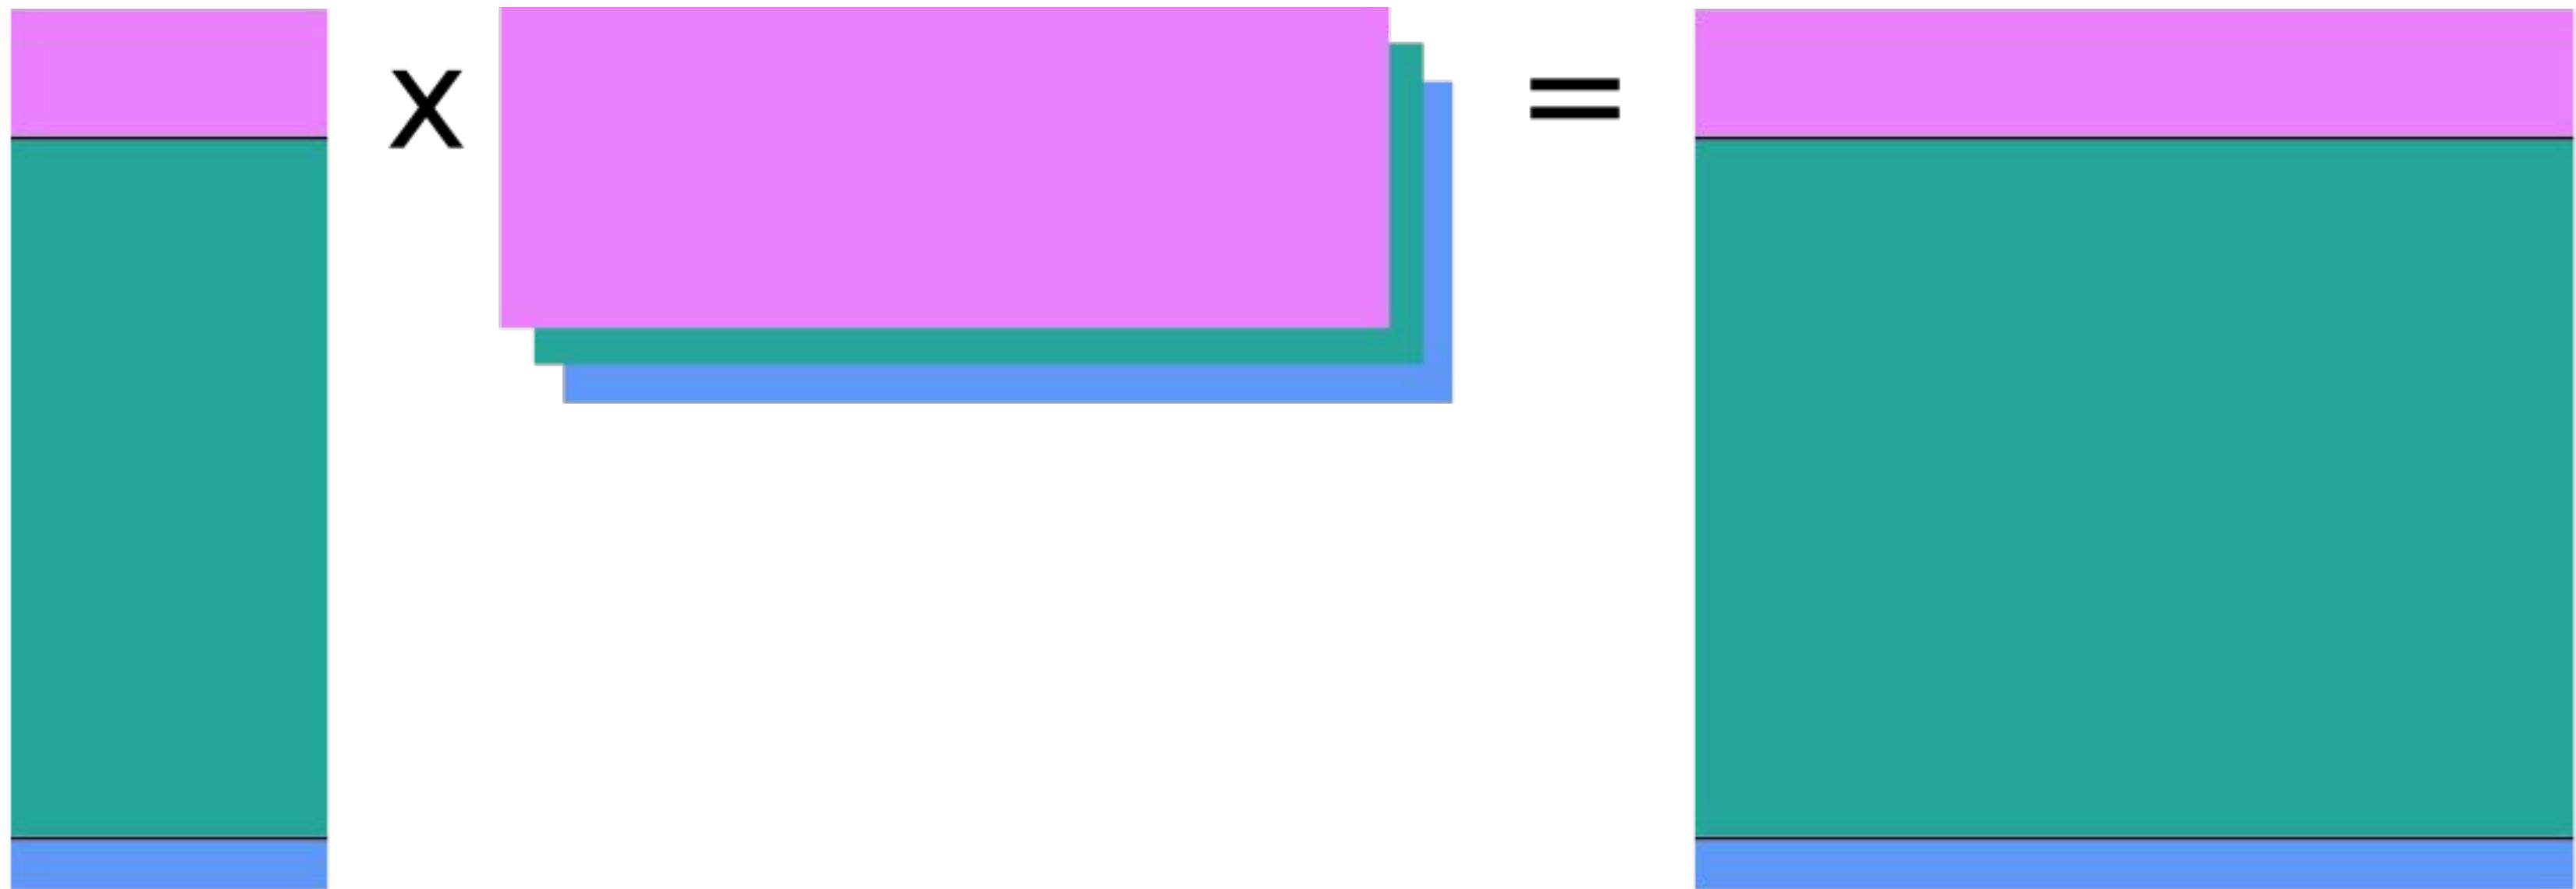
```
Compiling...: 100%|██████████| 64/64 [00:39<00:00, 1.62it/s]
Compiling...: 100%|██████████| 39/39 [00:02<00:00, 16.27it/s]
Profiling tpu:  0%|          | 0/5 [00:00<?, ?it/s]
Saving optimization profile to `/tmp/tuning_profile_2025-10-19_02:15:04_87s164fh`
```

id	block_m	block_k	block_n	t_mean (s)	t_std (s)
34	2048	512	2048	1.3220e-03	4.6820e-07
49	4096	512	1024	1.3353e-03	2.1772e-07
...					
4	512	1024	512	3.4113e-03	8.1699e-05
0	512	512	512	4.3446e-03	4.1993e-06

- worst result is 3.3x slower
- gap can get much wider for more complicated kernels
- tune_jax - available on PyPI
 - compiles in parallel (multi-core CPU speedup)
 - timing via automatic xprof parsing

future work & ragged dot on GPU

Future work



- collective fusions into the ragged dot kernel
- efficient quantization support
 - reusing quantized matrices in the backwards pass
 - in-kernel dynamic quantization

ragged dot on GPU

- originally in CUTLASS/CUBLAS
- kernel languages these days mostly
 - (many) triton implementations
 - e.g., (for JAX) https://github.com/rdyro/gpu_ragged_dot
 - not particularly efficient (no wmma)
- Mosaic GPU implementation
 - matmul itself more complicated (166 lines)
 - more dynamic control over tiles
 - github.com/jax-ml/jax...blackwell_ragged_dot_mgpu.py
 - super performant

Extras

jax-llm-examples

end-to-end inference

[jax-llm-examples](#) Public

main · 5 Branches · 0 Tags

Go to file · Add file · Code

rdro adding splash attention to qwen3 · ec8398a · 2 weeks ago · 86 Commits

- .github/workflows First draft of the GPT OSS model · 2 months ago
- deepseek_r1_jax serving cleanup · last month
- gpt_oss enable prefill splash attention in gpt_oss · last month
- kimi_k2 Implement ring-buffer attention for all models · 3 months ago
- llama3 serving cleanup · last month
- llama4 serving cleanup · last month
- misc Add TPU creation/deletion utilities; by gpolovets1 · 5 months ago
- qwen3 adding splash attention to qwen3 · 2 weeks ago
- serving enable prefill splash attention in gpt_oss · last month
- .gitignore serving cleanup · last month
- .pre-commit-config.yaml Add pre-commit, ruff linter and github CI · 4 months ago
- CONTRIBUTING.md Initial commit of repo basics · 8 months ago
- LICENSE Initial commit of repo basics · 8 months ago
- README.md First draft of the GPT OSS model · 2 months ago
- multi_host_README.md Add pre-commit, ruff linter and github CI · 4 months ago
- pyproject.toml Allow random quantized model initialization · 3 months ago

README · Contributing · Apache-2.0 license

JAX LLM examples

A collection (in progress) of example high-performance large language model implementations, written with JAX.

Current contents include:

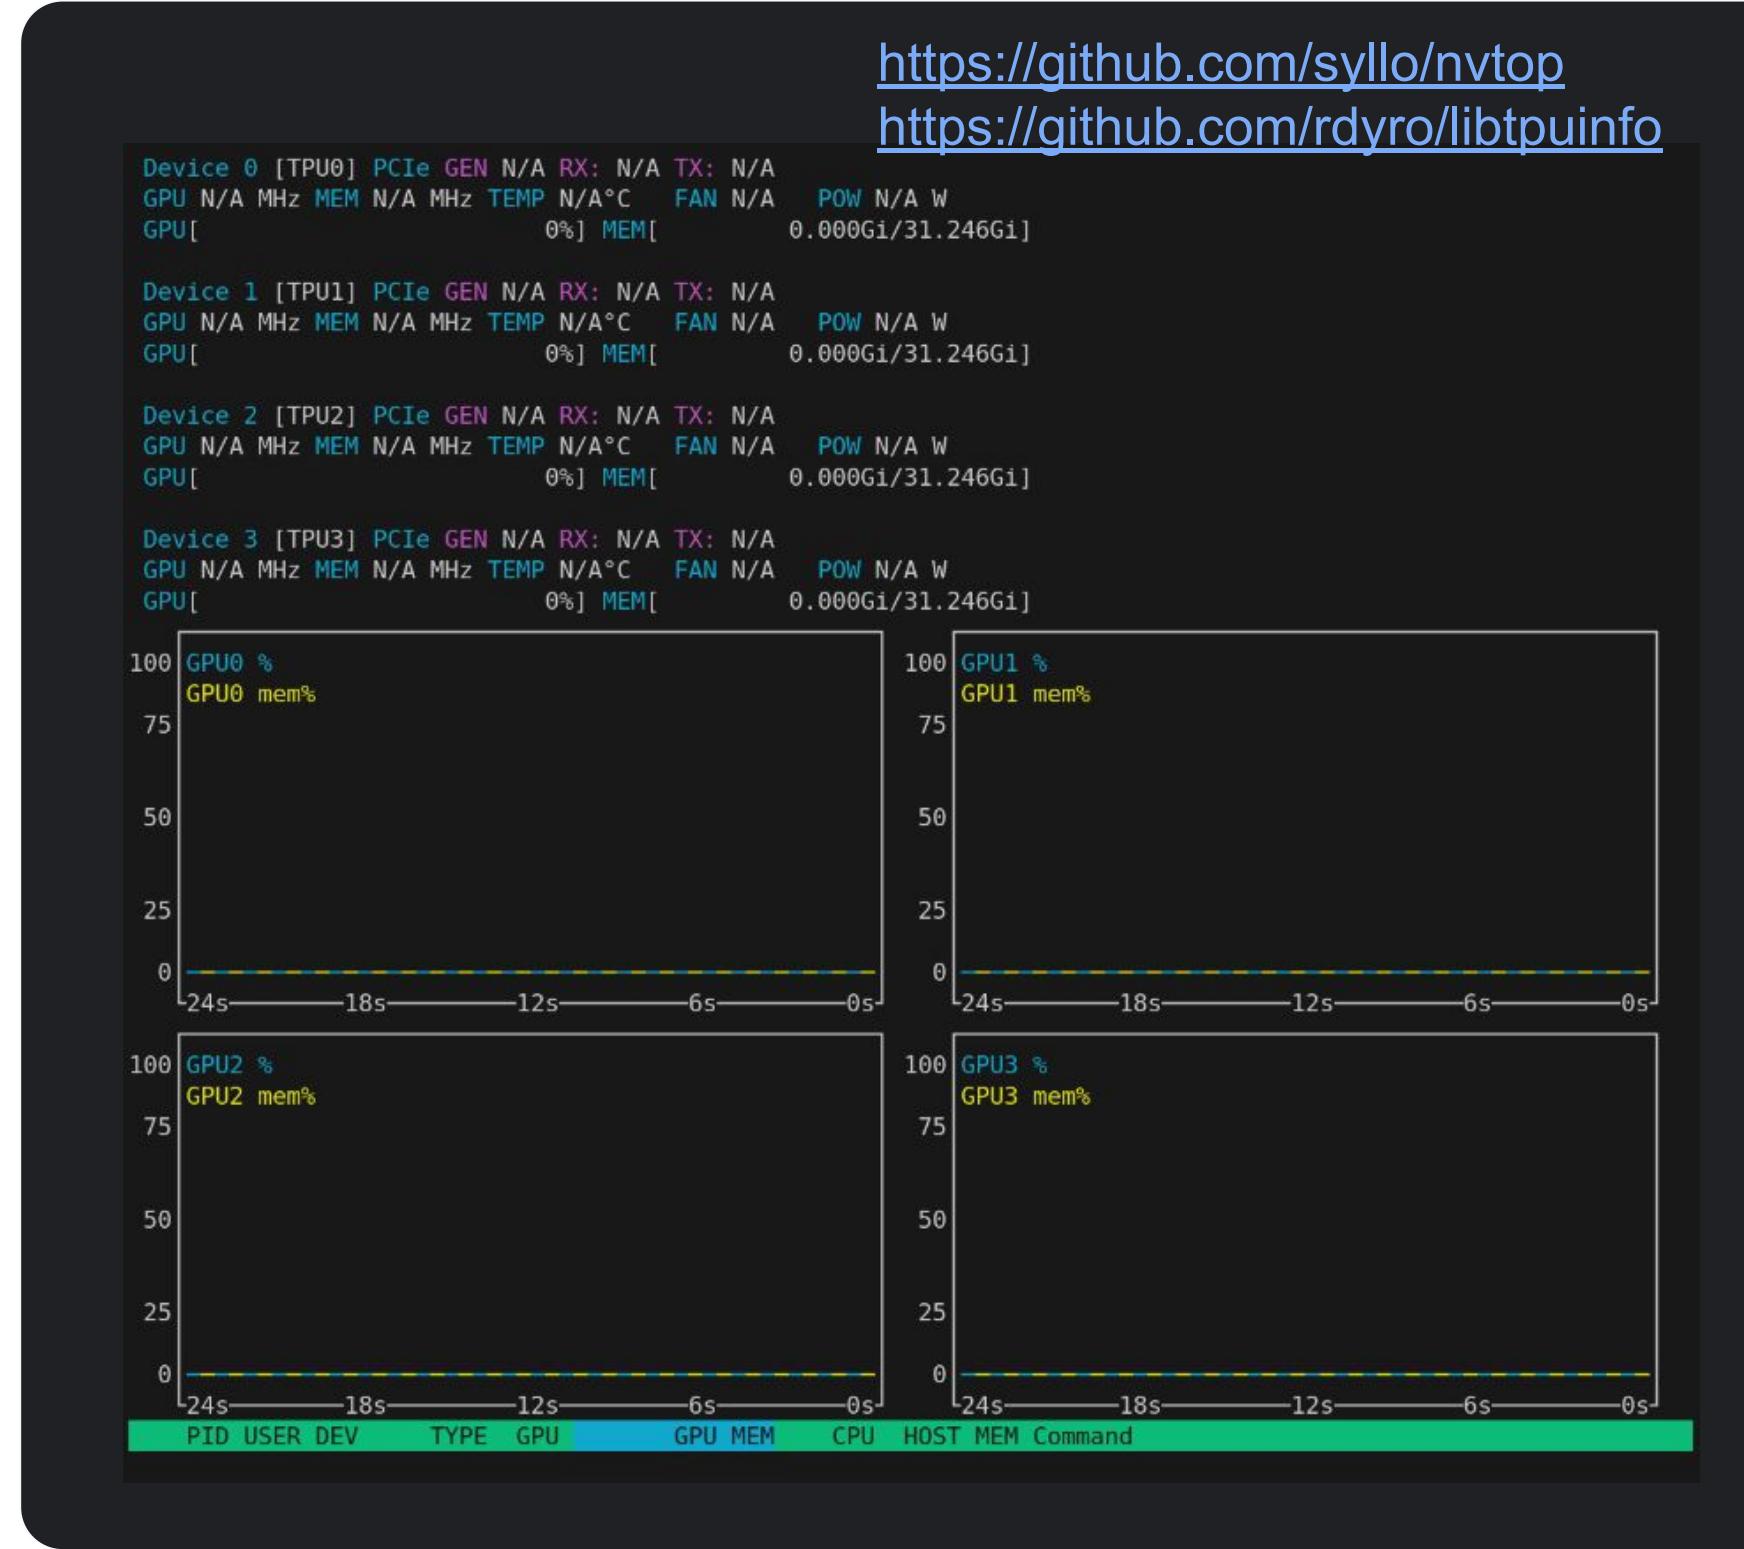
- DeepSeek R1
- Llama 4
- Llama 3
- Qwen 3
- Kimi K2
- OpenAI GPT OSS

For multi-host cluster setup and distributed training, see [multi_host_README.md](#) and the [tpu_toolkit.sh](#) script.

```
920     @partial(jax.shard_map, mesh=cfg.mesh, in_specs=in_specs, out_specs=out_spec, check_vma=False)
921     def _expert_fn(x, we_gate_up, we_gate_up_bias, we_down, we_down_bias, topk_weights, topk_idx):
922         (b, s, d), e = x.shape, cfg.moe_experts_per_tok
923         expert_idx = jax.lax.axis_index(expert_axname) if expert_axname is not None else 0
924         tensor_idx = jax.lax.axis_index(tensor_axname) if tensor_axname is not None else 0
925         topk_idx_ = topk_idx.reshape(-1)
926         valid_group_mask_ = (topk_idx_ >= expert_size * expert_idx) & (topk_idx_ < expert_size * (expert_idx + 1))
927         expert_mapped_topk_idx_ = jnp.where(valid_group_mask_, topk_idx_ - expert_idx * expert_size, 2**30)
928
929         sort_idx_ = jnp.argsort(expert_mapped_topk_idx_, axis=-1) # [b * s * e]
930         isort_idx_ = jnp.argsort(sort_idx_)
931
932         if cfg.ep_strategy == "prefill":
933             truncate_size = round(2 * sort_idx_.size / expert_count)
934             sort_idx_, isort_idx_ = sort_idx_[:truncate_size], isort_idx_[:truncate_size]
935
936             topk_idx_sort_ = topk_idx_[sort_idx_] # [b * s * e]
937             expert_mapped_topk_idx_sort_ = expert_mapped_topk_idx_[sort_idx_]
938             valid_group_mask_sort_ = expert_mapped_topk_idx_sort_ < 2**30
939             expert_mapped_topk_idx_sort_ = jnp.where(expert_mapped_topk_idx_sort_ < 2**30, expert_mapped_topk_idx_sort_, 0)
940
941             # equivalent to:
942             #
943             # x_repeat_ = jnp.repeat(x.reshape((-1, x.shape[-1])), e, axis=0)
944             # x_repeat_sort_ = jnp.take_along_axis(x_repeat_, sort_idx_[:, None], axis=-2) # [b * s, d]
945             #
946             x_repeat_sort_ = jnp.take_along_axis(x.reshape((-1, x.shape[-1])), sort_idx_[:, None] // e, axis=-2)
947             # [b * s * e, d] # "// e" is an index trick to avoid jnp.repeat
948
949             group_sizes = jnp.bincount(topk_idx_sort_, length=cfg.moe_numExperts)
950             group_sizes_shard = jax.lax.dynamic_slice_in_dim(group_sizes, expert_idx * expert_size, expert_size, 0)
951
952             with jax.named_scope("we_gate"):
953                 ff_gate_up = _moe_gmm(x_repeat_sort_, we_gate_up, group_sizes_shard, expert_mapped_topk_idx_sort_, cfg)
954                 ff_gate_up = ff_gate_up + we_gate_up_bias[expert_mapped_topk_idx_sort_, :]
955                 ff_gate = jnp.clip(ff_gate_up[..., ::2], max=cfg.moe_gateUpLimit)
956                 ff_up = jnp.clip(ff_gate_up[..., 1::2], min=-cfg.moe_gateUpLimit, max=cfg.moe_gateUpLimit)
957                 ff_gate_up = (ff_up + 1) * (ff_gate * jax.nn.sigmoid(ff_gate * cfg.moe_gateUpAlpha))
958                 ff_gate_up = jnp.where(valid_group_mask_sort_[:, None], ff_gate_up, 0)
959             with jax.named_scope("we_down"):
960                 ff_out = _moe_gmm(ff_gate_up, we_down, group_sizes_shard, expert_mapped_topk_idx_sort_, cfg)
961                 ff_out = ff_out + (tensor_idx == 0) * we_down_bias[expert_mapped_topk_idx_sort_, :]
962                 ff_out = jnp.where(valid_group_mask_sort_[:, None], ff_out, 0) # expensive
963
964             if cfg.ep_strategy == "prefill":
965                 rs_shape = math.ceil((ff_out.shape[-1] // tensor_count) / 256) * 256 * tensor_count
966                 pad_size = rs_shape - ff_out.shape[-1]
967                 ff_out = jnp.pad(ff_out, ((0, 0), (0, pad_size)))
968                 ff_out = jax.lax.psum_scatter(ff_out, axis_name=tensor_axname, scatter_dimension=1, tiled=True)
969                 ff_out = ff_out * topk_weights.reshape(1) # prevent divide by zero
```

Extras

nvtop: TPU support



```
 wget https://github.com/rdyro/libtpuinfo/releases/download/v0.0.1/libtpuinfo-linux-x86_64.so  
 sudo mv libtpuinfo-linux-x86_64.so /lib/libtpuinfo.so  
 git clone https://github.com/Syllo/nvtop.git  
 cd nvtop && mkdir build && cd build && cmake -DTPU_SUPPORT=ON .. && make  
 sudo cp ./src/nvtop /usr/local/bin
```